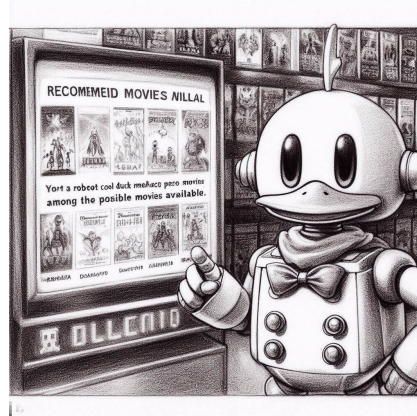


Introduction to Natural Language Processing



AI frameworks

Course: 8 videos

- Introduction
- Tokenization
- Bag-of-words
- Co-occurrence Matrix
- Word embeddings
- Sequential modeling
- Transformers
- Large language models



Introduction



What is a language

A language is a structured system of communication.

The structure of a language is its grammar and the free components are its vocabulary.

Languages are the primary means of communication of humans, and can be conveyed through speech (spoken language), sign, or writing.

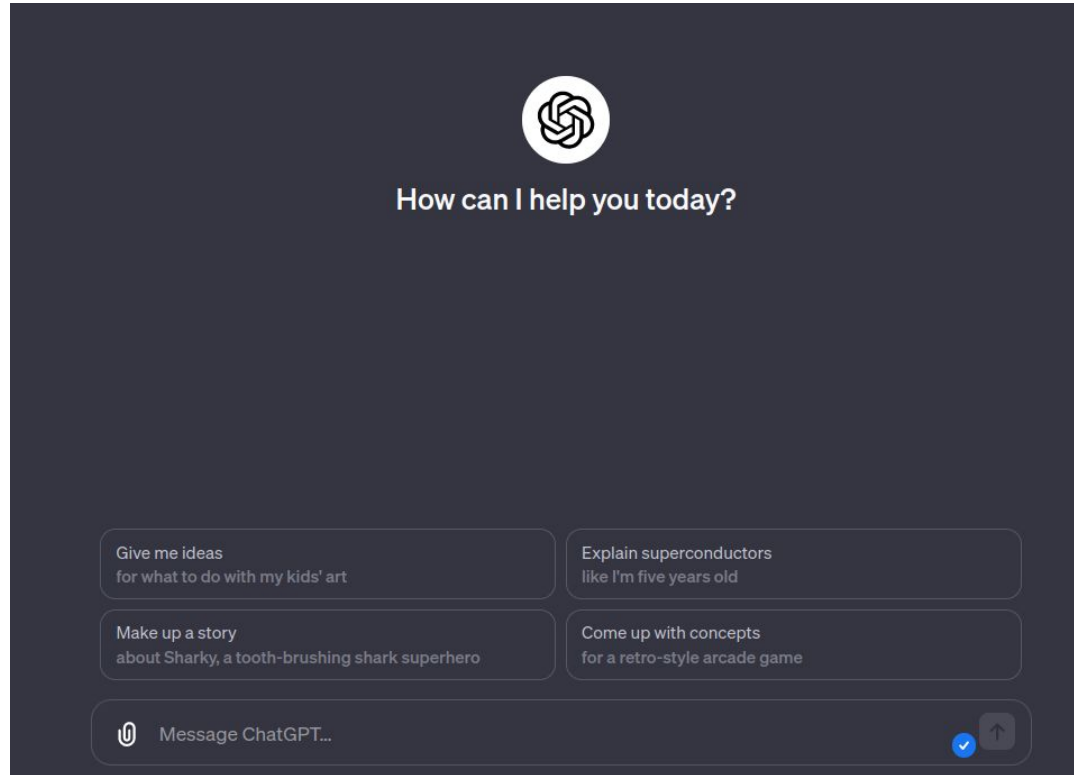
(Source Wikipedia)

Natural Language Processing

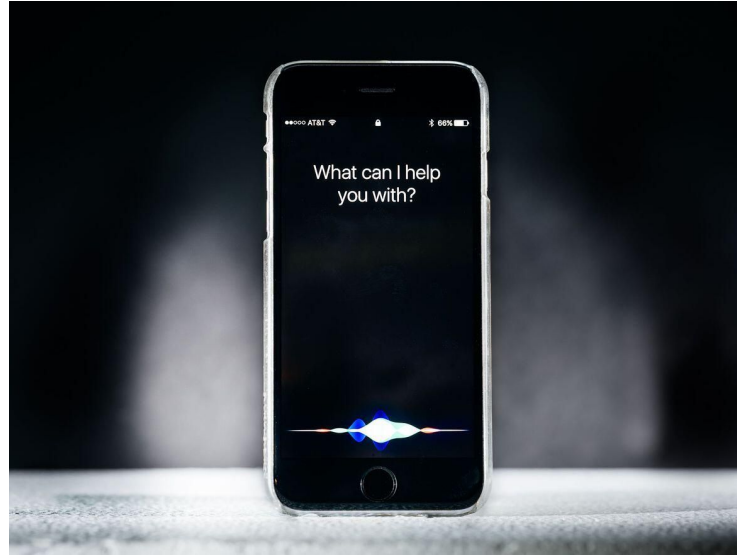
Objectives:

- Design programs able to understand human language as it is spoken and written
- Extract insightful information
- Generate language

Examples



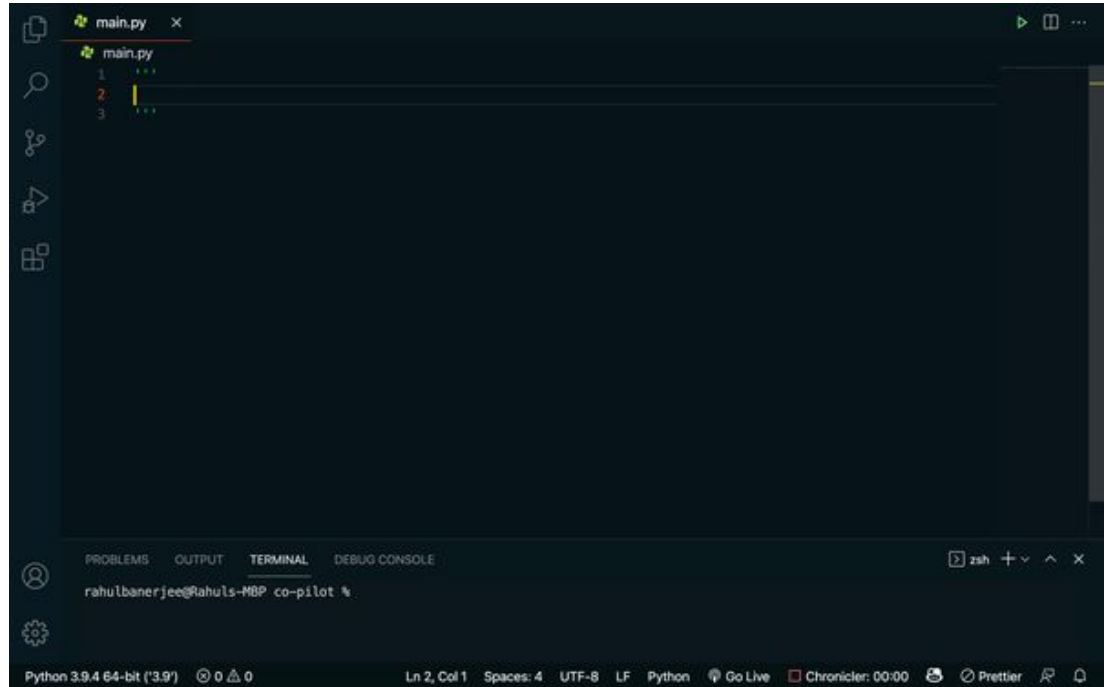
Examples



Examples



GitHub
Copilot



Examples



Captioning Model

A happy dog is standing in the ocean

Example

 DALL·E 2

“A manga-style drawing of the cool duck listening to music while coding.”



NLP tasks

(X, Y) random variables

Estimate the conditional probability $P(Y|X=x)$

- **Word:** *chaise*
- **Word sequence:** [*il, porte, une, chemise*]
- **Document:** *Il porte une chemise jaune. Ses chaussures sont noires ...*

NLP tasks

(X,Y) random variables

Estimate the conditional probability $P(Y|X=x)$

- **Word:** *chaise*
- **Word sequence:** *[il, porte, une, chemise]*
- **Document:** *Il porte une chemise jaune. Ses chaussures sont noires ...*
- Classification

NLP tasks

(X, Y) random variables

Estimate the conditional probability $P(Y|X=x)$

- **Word:** *chaise*
- **Word sequence:** [*il, porte, une, chemise*]
- **Document:** *Il porte une chemise jaune. Ses chaussures sont noires ...*
- Classification
- Translation, Text generation, Question answering, Text summarization, ...
- Document ranking, recommendations, ...

Representations

Machine learning algorithms need vectorial representations

- Words -> vectors: Voiture -> (0.2, 4.1, ..., -2.3,)
- Documents -> vectors La voiture roule -> (3.3, -2.1, ..., 2.5)

We call these vector representation **embeddings**

Two approaches:

- Statistical representations -> Bag of Words
- Learned representations -> Word embeddings

Textual data

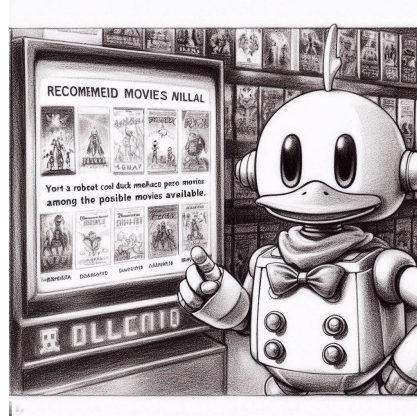
- A **corpus** is a collection of:
 - **documents** which are sequences of :
 - **tokens**

Token: basic unit of discrete data indexed from a vocabulary

- Word
- Sub-word
- A sequence of words or sub-words
- A character
- A symbol

How to identify tokens?

Introduction to Natural Language Processing



AI frameworks



Tokenization



Textual data

- A **corpus** is a collection of:
 - **documents** which are sequences of :
 - **tokens**

Token: basic unit of discrete data indexed from a vocabulary

- Word
- Sub-word
- A sequence of words or sub-words
- A character
- A symbol

How to identify tokens?

Tokenization

Consists in segmenting a document into tokens

Not so trivial!

This is an example of tokenization.

Use whitespace?

[“This”, “is”, “an”, “example”, “of”, “tokenization.”]

Tokenization

Consists in segmenting a document into tokens

Not so trivial!

This is an example of tokenization.

Use whitespace and punctuation?

["This", "is", "an", "example", "of", "tokenization", "."]

That's a problem...

["That", "'", "s", "a", "problem", ".", ".", "."]

Tokenization

Many other problems:

- Compound words (e.g. pick-pocket, German, ...)
- No separators (Chinese, Japanese)
- Casual language: This is a coool #dummysmiley: :-) :-P <3
- ...

Many partial solutions:

- Character level tokenization
- Regular Expression tokenization
- Dictionary based tokenization
- Rule Based Tokenization (Penn TreeBank, Spacy, Moses, TweetTokenizer ...)

Lemmatization

Lemmatization is the process of grouping inflected forms together as a single base form

Example:

"builds", "building", or "built" => "build"

Stemming

Stemming is the process of reducing inflected words to their word stem, base or root form

Example:

“programming”, “programs”, “programmed” => “program”

Information loss

Stemming and Lemmatization may lead to crucial information loss

Example: sentiment analysis

objective, objection -> object

compliment, complicate -> comply

Subwords tokenization

Principle:

- Frequently used words are unit tokens
- Less frequent words should be decomposed into meaningful subwords
e.g. “annoyingly” => “annoying” and “ly”
- Rely on model training to discover the most frequent occurring pairs of symbols

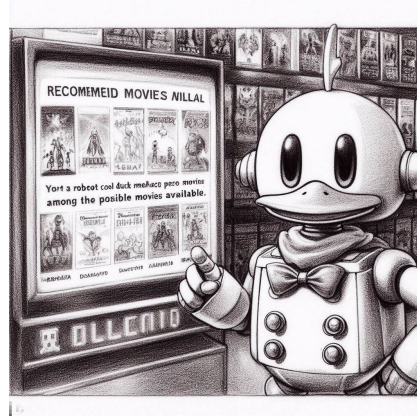
Advantages:

- Reasonable vocabulary size
- Process unknown words

Text Cleaning

- Remove noise (**HTML, punctuation, specific symbols like #, nouns, references, ...**)
- Remove stop words (**and, or, the, ...**)
- Pass a spell checker on your data
- Convert to lowercase

Introduction to Natural Language Processing



AI frameworks



Bag of Words



Vectorizing documents



[0.2, 5.1, ..., -4.1, 2.2]
[-0.8, 2.1, ..., 2.1, 0.7]
[2.1, -3.3, ..., 0.8, 2.5]

Principle :

Use tokens statistics

One hot encoding

- (1) I am going to the supermarket
- (2) The post office is close to the supermarket

I	am	going	to	the	supermarket	post	office	is	close
1	1	1	1	1	1	0	0	0	0
0	0	0	1	1	0	1	1	1	1

Term-frequency

- (1) I am going to the supermarket
(2) The post office is close to the supermarket

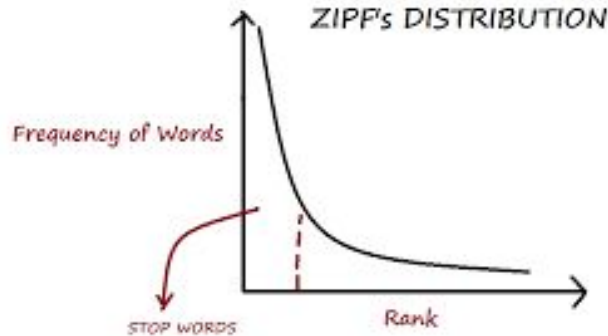
I	am	going	to	the	supermarket	post	office	is	close
1	1	1	1	1	1	0	0	0	0
0	0	0	1	2	0	1	1	1	1

Term-frequency matrix: $tf_{t,d} = |\{t \in d\}|$

TF-IDF

Count based encoding is sensitive to frequent words

- A word present in all documents **is not very informative**
- A word present in few documents **is very informative**



Weight term frequency with the **Inverse Document Frequency**:

$$idf_{t,C} = \log \left(\frac{|C|}{|\{d \in C, s.t. t \in d\}|} \right)$$

TF-IDF

(1) This is a big supermarket

(2) This post office is close to a supermarket

This	is	a	to	supermarket	post	office	big	close
0	0	0	0	0	0	0	0.3	0
0	0	0	0.3	0	0.3	0.3	0	0.3

$$tf_{t,d} \cdot idf_{t,C} = |\{t \in d\}| \log \left(\frac{|C|}{|\{d \in C, s.t. t \in d\}|} \right)$$

Bag of words

How to deal with negation or context:

(1) **I do not like carrots**

(2) **I do like carrots**

I	do	not	like	carrots
1	1	1	1	1
1	1	0	1	1

Bag of words

How to deal with negation or context:

(1) **I do not like carrots**

(2) **I do like carrots**

I	do	not	like	carrots
1	1	1	1	1
1	1	0	1	1

Using n-grams

I	I do	do not	not like	like carrots	carrots
1	1	1	1	1	1

Bag of words pipeline:

<body>I was thinking of eating all the fruits :-)</body>

1. Text cleaning, removing noise, lower case:
i was thinking of eating all the fruits
2. Remove stop words:
i was thinking eating all fruits
3. Stemming or lemmatization:
i was think eat fruit
4. Tokenize:
[i, was, think, eat, fruit]

Bag of words pipeline:

<body>I was thinking of eating all the fruits :-)</body>

1. Text cleaning, removing noise, lower case:
i was thinking of eating all the fruits
2. Remove stop words:
i was thinking eating all fruits
3. Stemming or lemmatization:
i was think eat fruit
4. Tokenize:
[i, was, think, eat, fruit]
5. Vectorize (One-hot encoding, count, tf-idf ...)

I	potato	think	drink	...	fruit
0.1	0	0.5	0	...	0.73
0	0.11	0.4	0.22	...	0.1

Bag of words pipeline:

<body>I was thinking of eating all the fruits :-)</body>

1. Text cleaning, removing noise, lower case:
i was thinking of eating all the fruits
2. Remove stop words:
i was thinking eating all fruits
3. Stemming or lemmatization:
i was think eat fruit
4. Tokenize:
[i, was, think, eat, fruit]
5. Vectorize (One-hot encoding, count, tf-idf ...)

I	potato	think	drink	...	fruit
0.1	0	0.5	0	...	0.73
0	0.11	0.4	0.22	...	0.1

6. Apply any ML algorithm

Other features

- Number of words, characters, ...
- Grammatical categories
- Number of capital letters
- Mean word length
- ...

Bag of words

Advantages:

- Easy to set up
- No training required

Drawbacks

- Does not scale with vocabulary size
- Very sparse representation
- No semantic information (synonyms are treated like different words)

Bag of words

Advantages:

- Easy to set up
- No training required

Drawbacks

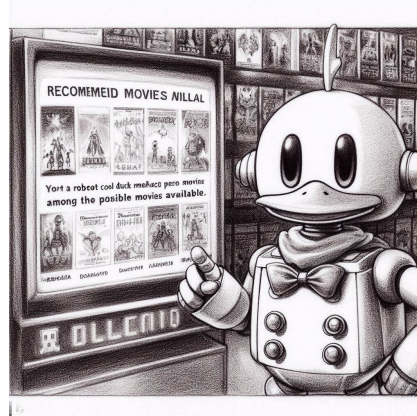
- Does not scale with vocabulary size
- Very sparse representation
- No semantic information (synonyms are treated like different words)

1/ I like this movie

2/ I love this film

I	like	this	movie	love	that	film
1	1	1	1	0	0	0
1	0	0	0	1	1	1

Introduction to Natural Language Processing



AI frameworks



Co-occurrence Matrix



Vectorizing documents



[0.2, 5.1, ..., -4.1, 2.2]

[-0.8, 2.1, ..., 2.1, 0.7]

[2.1, -3.3, ..., 0.8, 2.5]

Principle :

Use tokens statistics

Vectorizing words

car
house
fruit



[0.2, 5.1, ..., -4.1, 2.2]

[-0.8, 2.1, ..., 2.1, 0.7]

[2.1, -3.3, ..., 0.8, 2.5]

Principle :

Use words contexts

A word is defined by its context

Tous les matins je mets du “___” dans mes céréales

-> lait, sucre, ...

Tu joues toujours du Mayuri?



A word is defined by its context

Tous les matins je mets du “___” dans mes céréales

-> lait, sucre, ...

Tu joues toujours du Mayuri?

Two main approaches:

- Frequency-based
- Prediction-based



A word is defined by its context

Tous les matins je mets du “___” dans mes céréales

-> lait, sucre, ...

Tu joues toujours du Mayuri?

Two main approaches:

- Frequency-based
- Prediction-based



Co-occurrence Matrix

- (1) For breakfast I eat eggs and fruits
- (2) Go buy some eggs at the supermarket
- (3) There were no eggs left at the supermarket

	egg	fruit	...	supermarket
egg	-	1	..	2
fruit	1	-	...	0
...
supermarket	2	0	...	-

Co-occurrence Matrix

- (1) For breakfast I eat eggs and fruits
- (2) Go buy some eggs at the supermarket
- (3) There were no eggs left at the supermarket

	egg	fruit	...	supermarket
egg	-	1	..	2
fruit	1	-	...	0
...
supermarket	2	0	...	-

Co-occurrence Matrix

- Embedding magnitudes depend on the corpus size
- Frequent words have a strong impact

Co-occurrence Matrix

- Embedding magnitudes depend on the corpus size
- Frequent words have a strong impact

-> Pointwise Mutual Information (PMI)

$$\text{PMI}(w_1, w_2) = \log \frac{p(w_1, w_2)}{p(w_1)p(w_2)}$$

Diagram illustrating the components of the PMI formula:

- $p(w_1, w_2)$: Probability of w_1 and w_2 co-occurring (represented by the brown box in the numerator).
- $p(w_1)$: Probability of w_1 occurring (represented by the blue box in the denominator).
- $p(w_2)$: Probability of w_2 occurring (represented by the yellow box in the denominator).

Co-occurrence Matrix

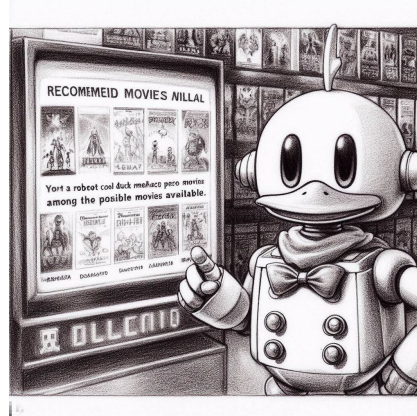
$$PMI(w_1, w_2) = \log \frac{\frac{1}{n_{pairs}} \# \{(w_1, w_2)\}}{\frac{1}{n_{word}} \# \{w_1\} \frac{1}{n_{word}} \# \{w_2\}}$$

- Embedding magnitudes depend on the corpus size
- Frequent words have a strong impact

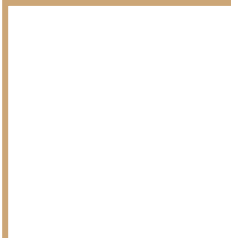
-> **Pointwise Mutual Information (PMI)**

- Large embedding space **$O(nb_words^2)$**
- Singular Value Decomposition (SVD)


Introduction to Natural Language Processing



AI frameworks



Word embeddings



Vectorizing words

car
house
fruit



[0.2, 5.1, ..., -4.1, 2.2]

[-0.8, 2.1, ..., 2.1, 0.7]

[2.1, -3.3, ..., 0.8, 2.5]

Principle :

Use words contexts

A word is defined by its context

Tous les matins je mets du “___” dans mes céréales

-> lait, sucre, ...

Tu joues toujours du Mayuri?

Two main approaches:

- Frequency-based
- Prediction-based



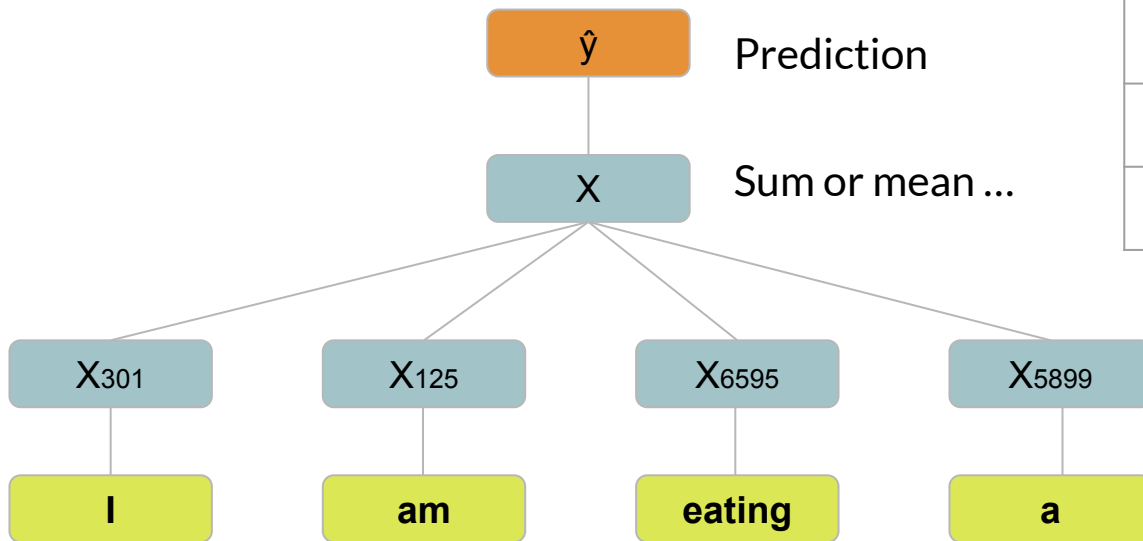
Word2Vec

- Words are mapped to embeddings (e.g. **eating** -> [3.1, -2.5,...,1.8])

Word	idx	Embeddings
I	X ₃₀₁	[0.2, 1.1,...,-1.2]
...	...	
am	X ₁₂₅	[-1.5, 2.3,...,0.2]
...	...	
eating	X ₆₅₉₅	[3.1, -2.5,...,1.8]

Word2Vec

Embeddings are built through shallow neural networks trained to reconstruct linguistic contexts of words



Word	idx	Embeddings
I	X_{301}	[0.2, 1.1,...,-1.2]
...	...	
am	X_{125}	[-1.5, 2.3,...,0.2]
...	...	
eating	X_{6595}	[3.1, -2.5,...,1.8]

Embedding look-up

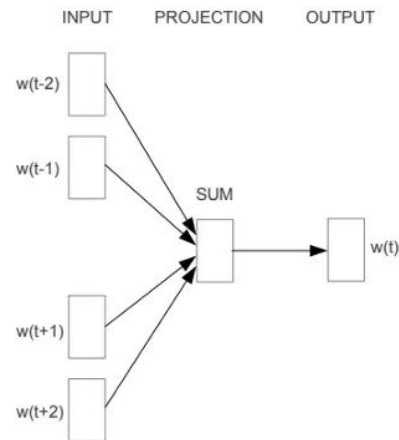
Word2Vec

Source Texte

The quick brown fox jumps over the lazy dog.

The quick brown fox jumps over the lazy dog.

The quick brown fox jumps over the lazy dog.



CBOW

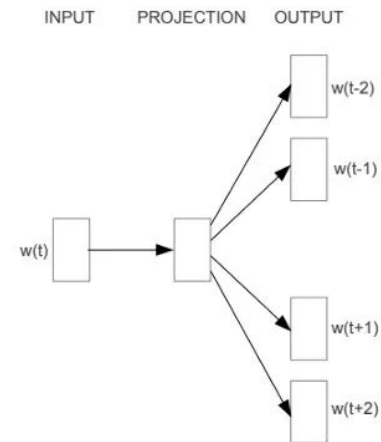
Word2Vec

Source Texte

The quick brown fox jumps over the lazy dog.

The quick brown fox jumps over the lazy dog.

The quick brown fox jumps over the lazy dog.



Skip-gram

Negative sampling

- Predict word given a context is a classification task

Softmax:
$$p(w_O \mid c_I) = \frac{\exp(v'_{w_O}{}^T v_{c_I})}{\sum_{w=1}^W \exp(v'_w{}^T v_{c_I})}$$

Negative sampling

- Predict word given a context is a classification task

$$\text{Softmax: } p(w_O \mid c_I) = \frac{\exp(v'_{w_O}{}^T v_{c_I})}{\sum_{w=1}^W \exp(v'_w{}^T v_{c_I})}$$

- Denominator is expensive to compute

Negative sampling

- Predict word given a context is a classification task

Softmax: $p(w_O \mid c_I) = \frac{\exp(v'_{w_O}{}^T v_{c_I})}{\sum_{w=1}^W \exp(v'_w{}^T v_{c_I})}$

- Denominator is expensive to compute

Solution: **Negative Sampling:**

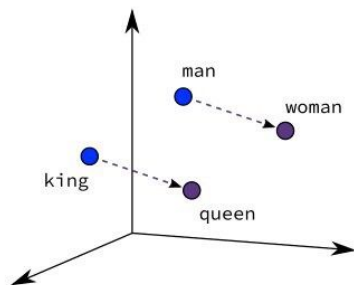
- Maximize the target probability

- Sample k negative samples and minimize their probability

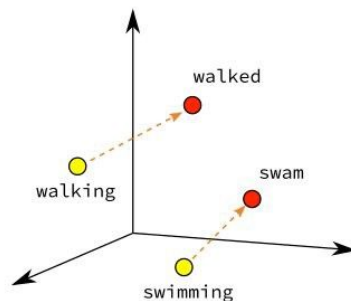
$$\underbrace{\log \sigma(v'_{w_O}{}^T v_{w_I})}_{\text{Maximize}} + \sum_{i=1}^k \underbrace{\mathbb{E}_{w_i \sim P_n(w)} \left[\log \sigma(-v'_{w_i}{}^T v_{w_I}) \right]}_{\text{Minimize}}$$

Word2Vec

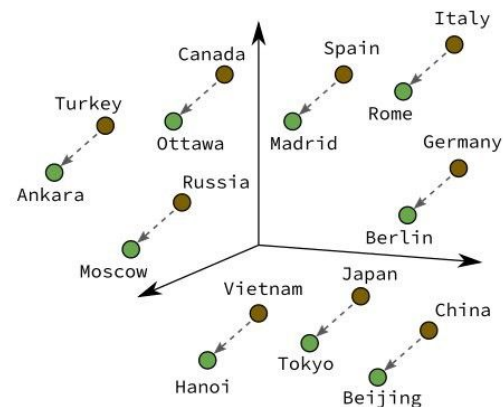
Word with similar **semantic** are located close to one another in the **latent space**



Male-Female



Verb Tense



Country-Capital

Word2Vec: Analogies

king - man + woman \approx queen



How to get document embedding?

Word	idx	Embeddings
the	X ₃₀₁	[0.2, 1.1,...,-1.2]
...	...	
baby	X ₁₂₅	[-1.5, 2.3,...,0.2]
...	...	
is	X ₆₅₉₅	[3.1, -2.5,...,1.8]

X₃₀₁

the

X₁₂₅

baby

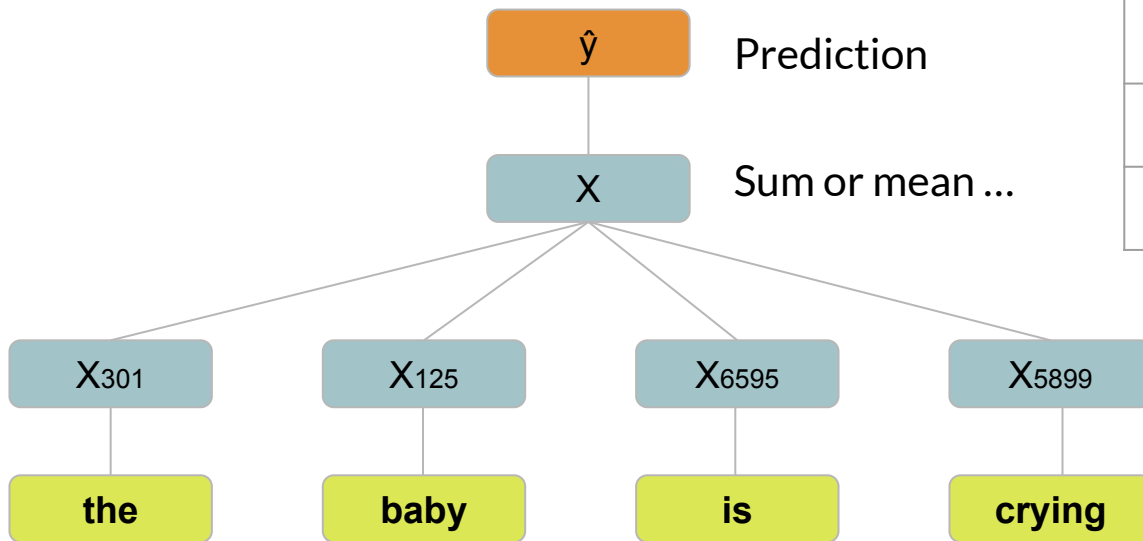
X₆₅₉₅

is

X₅₈₉₉

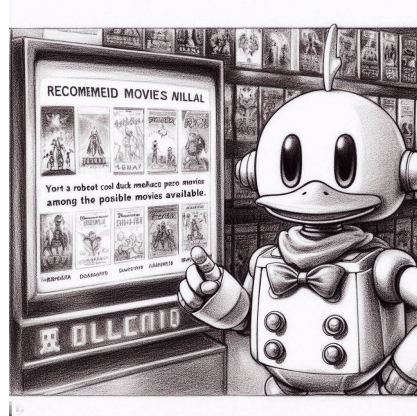
crying

How to get document embedding?

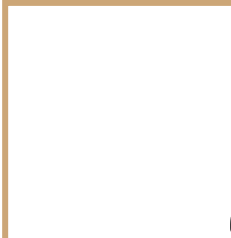


Word	idx	Embeddings
the	X_{301}	[0.2, 1.1,...,-1.2]
...	...	
baby	X_{125}	[-1.5, 2.3,...,0.2]
...	...	
is	X_{6595}	[3.1, -2.5,...,1.8]


Introduction to Natural Language Processing

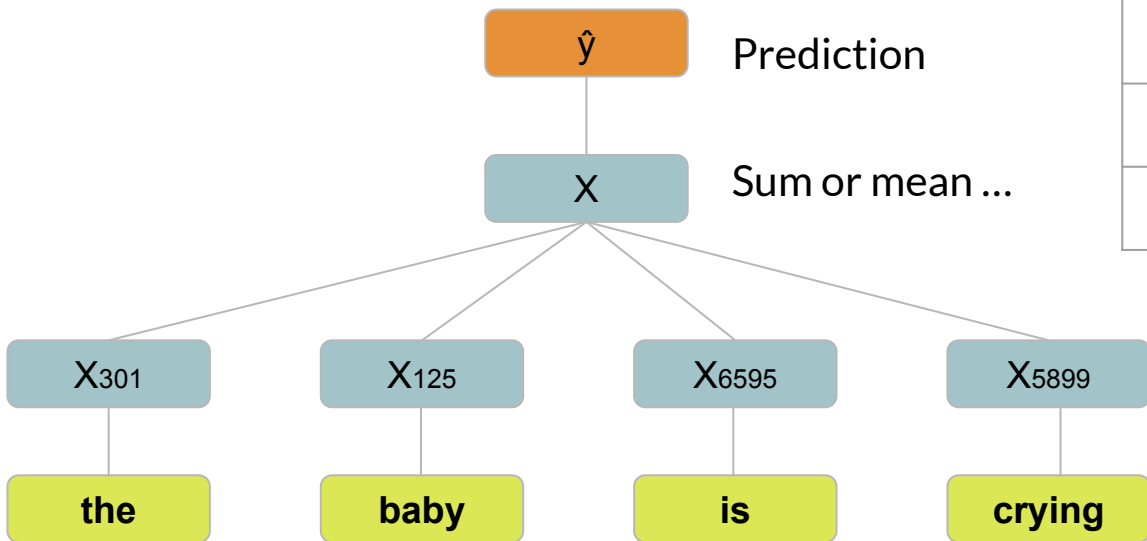


AI frameworks



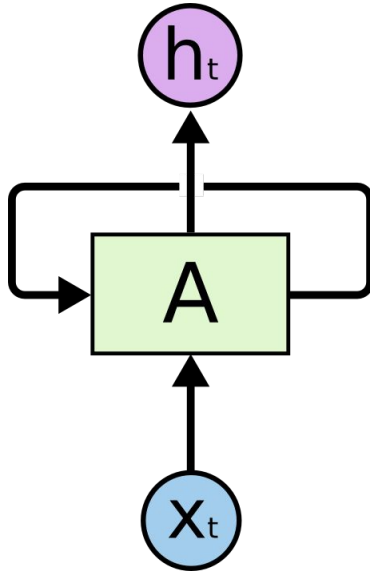
Sequence modeling



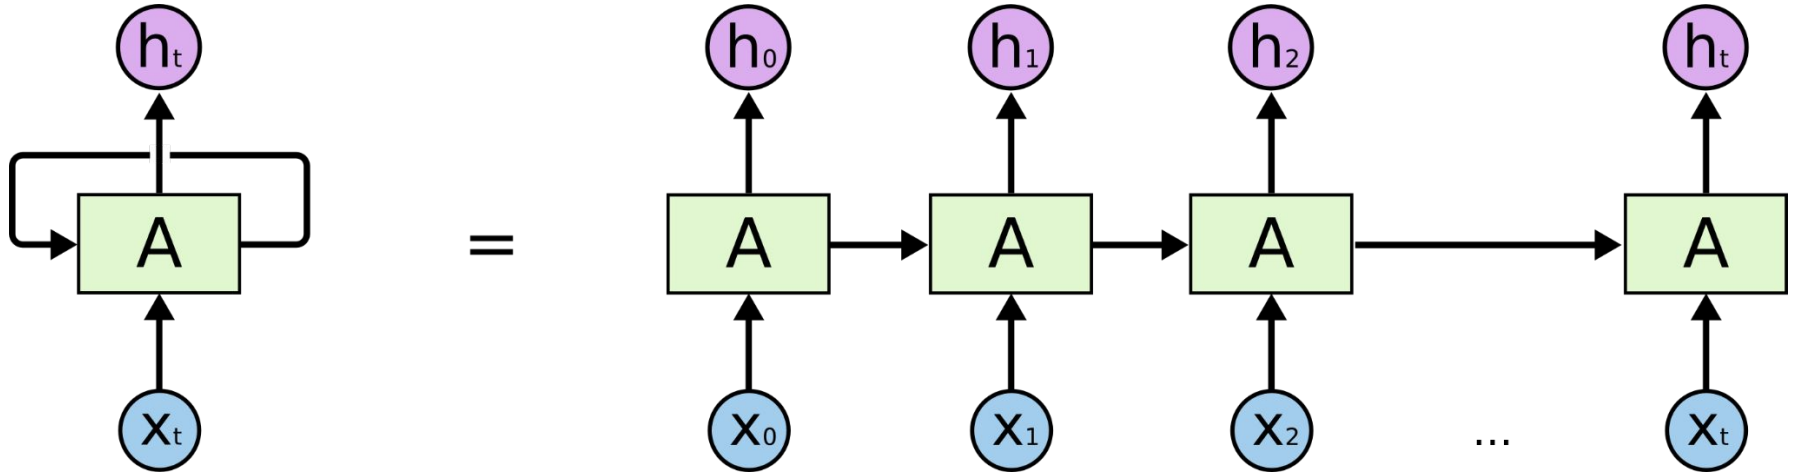


Word	idx	Embeddings
the	X_{301}	$[0.2, 1.1, \dots, -1.2]$
...	...	
baby	X_{125}	$[-1.5, 2.3, \dots, 0.2]$
...	...	
is	X_{6595}	$[3.1, -2.5, \dots, 1.8]$

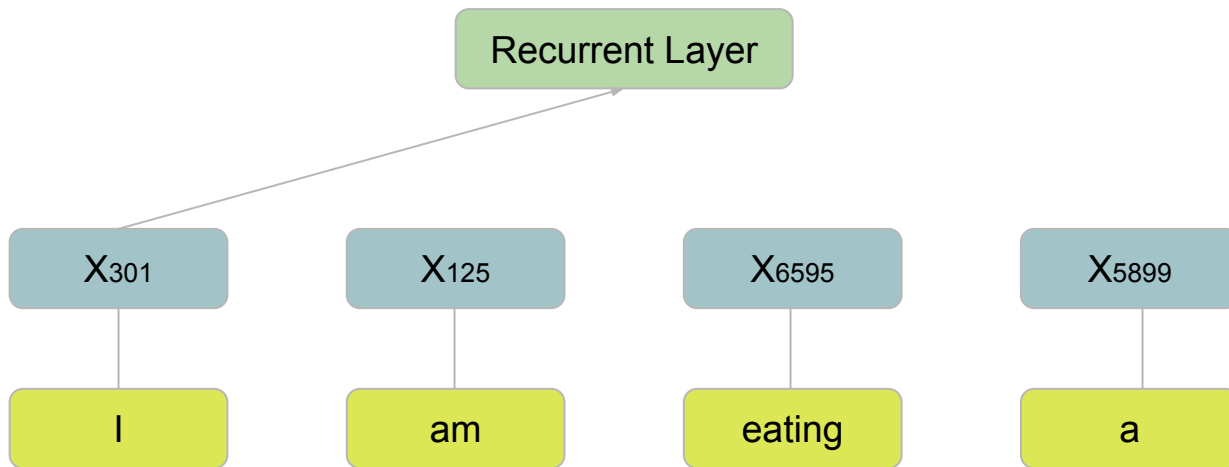
Recurrent layer



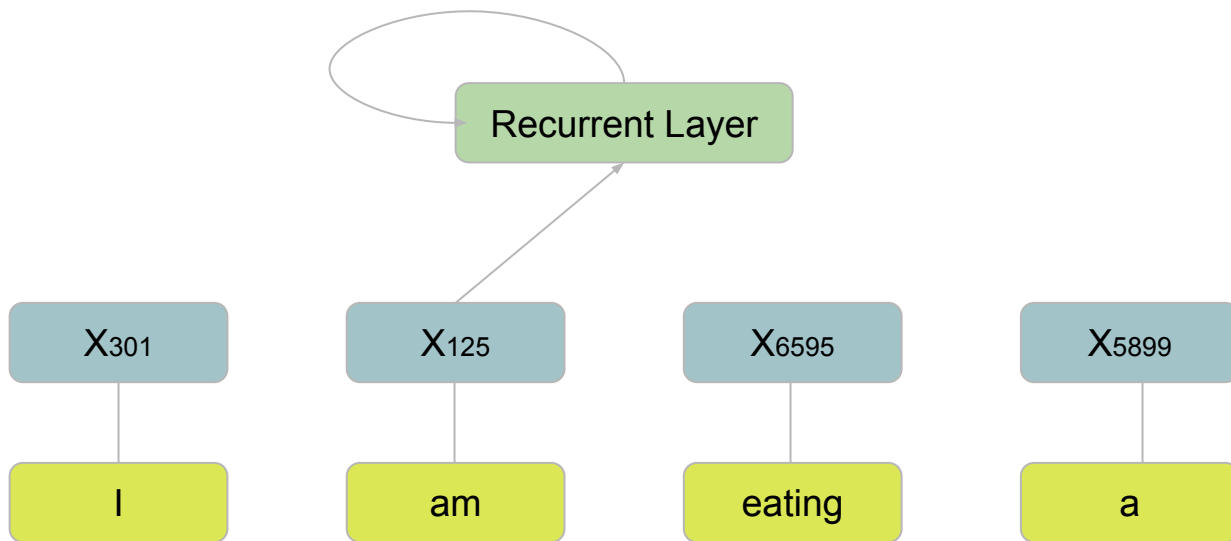
Recurrent layer



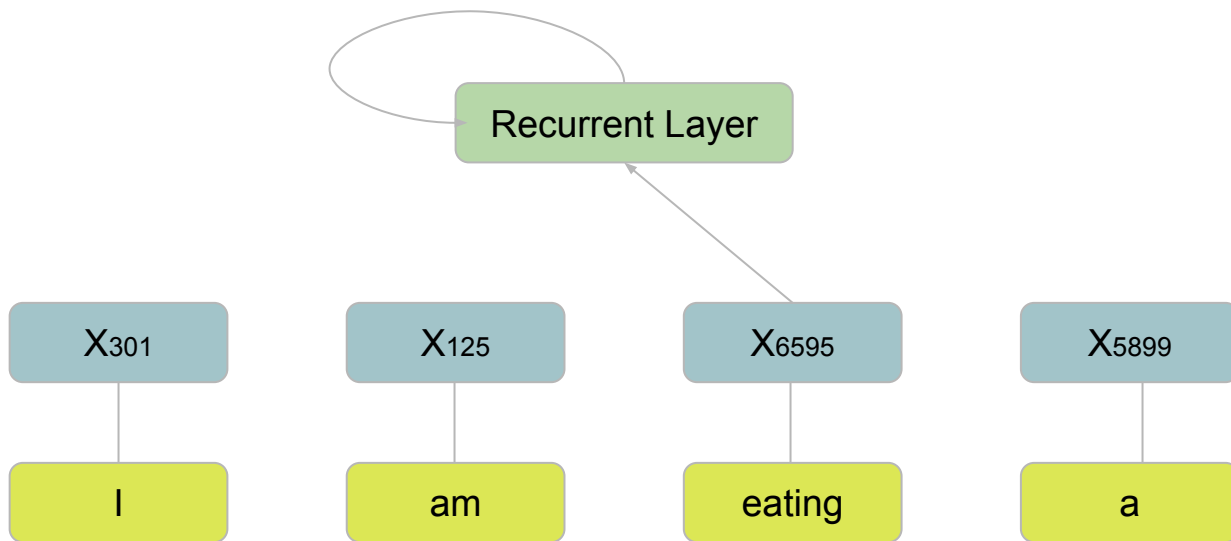
Sequence modeling



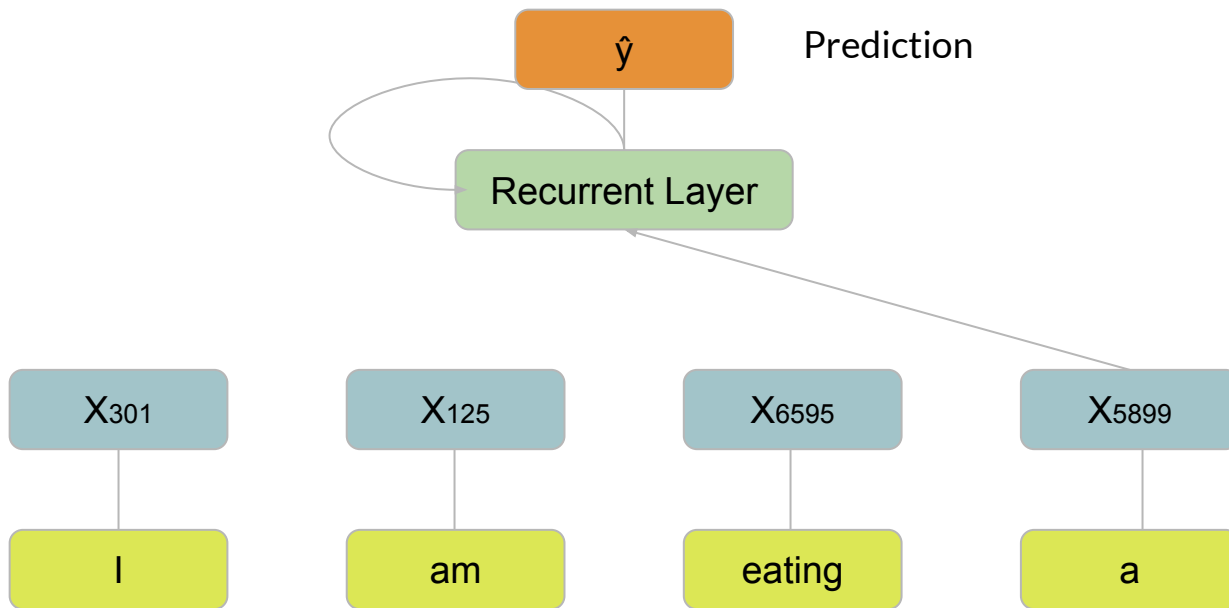
Sequence modeling



Sequence modeling

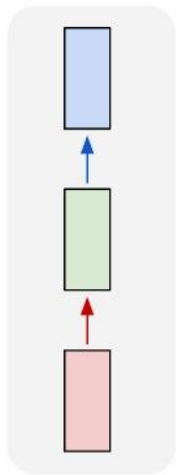


Sequence modeling

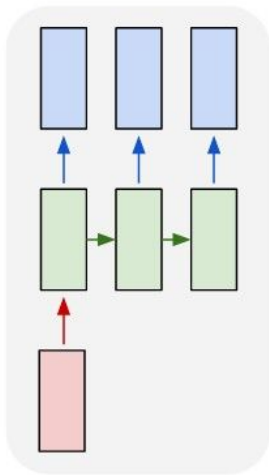


Sequence modeling

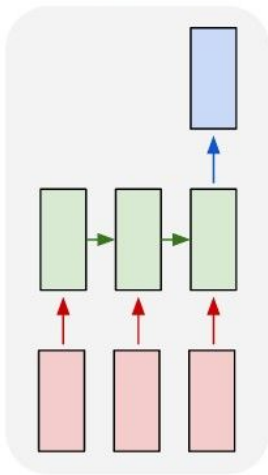
one to one



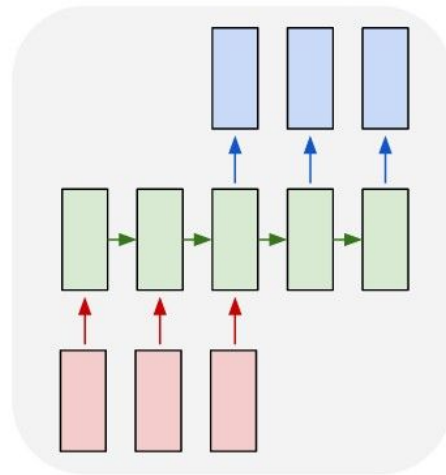
one to many



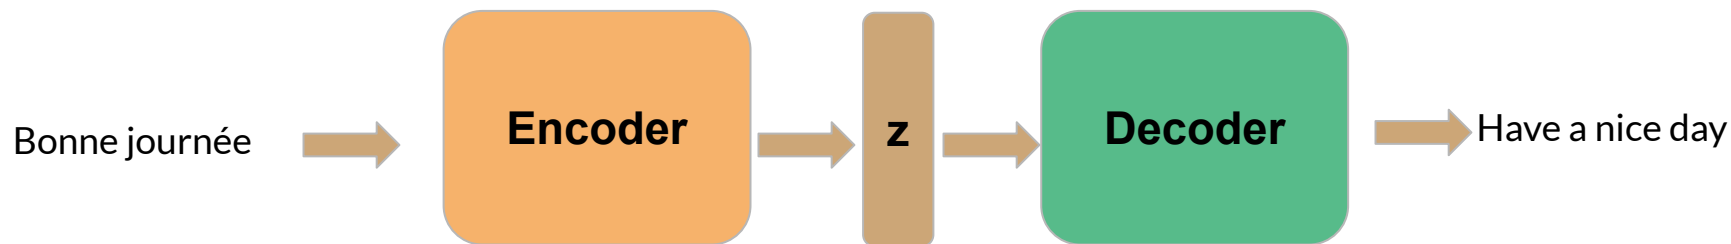
many to one



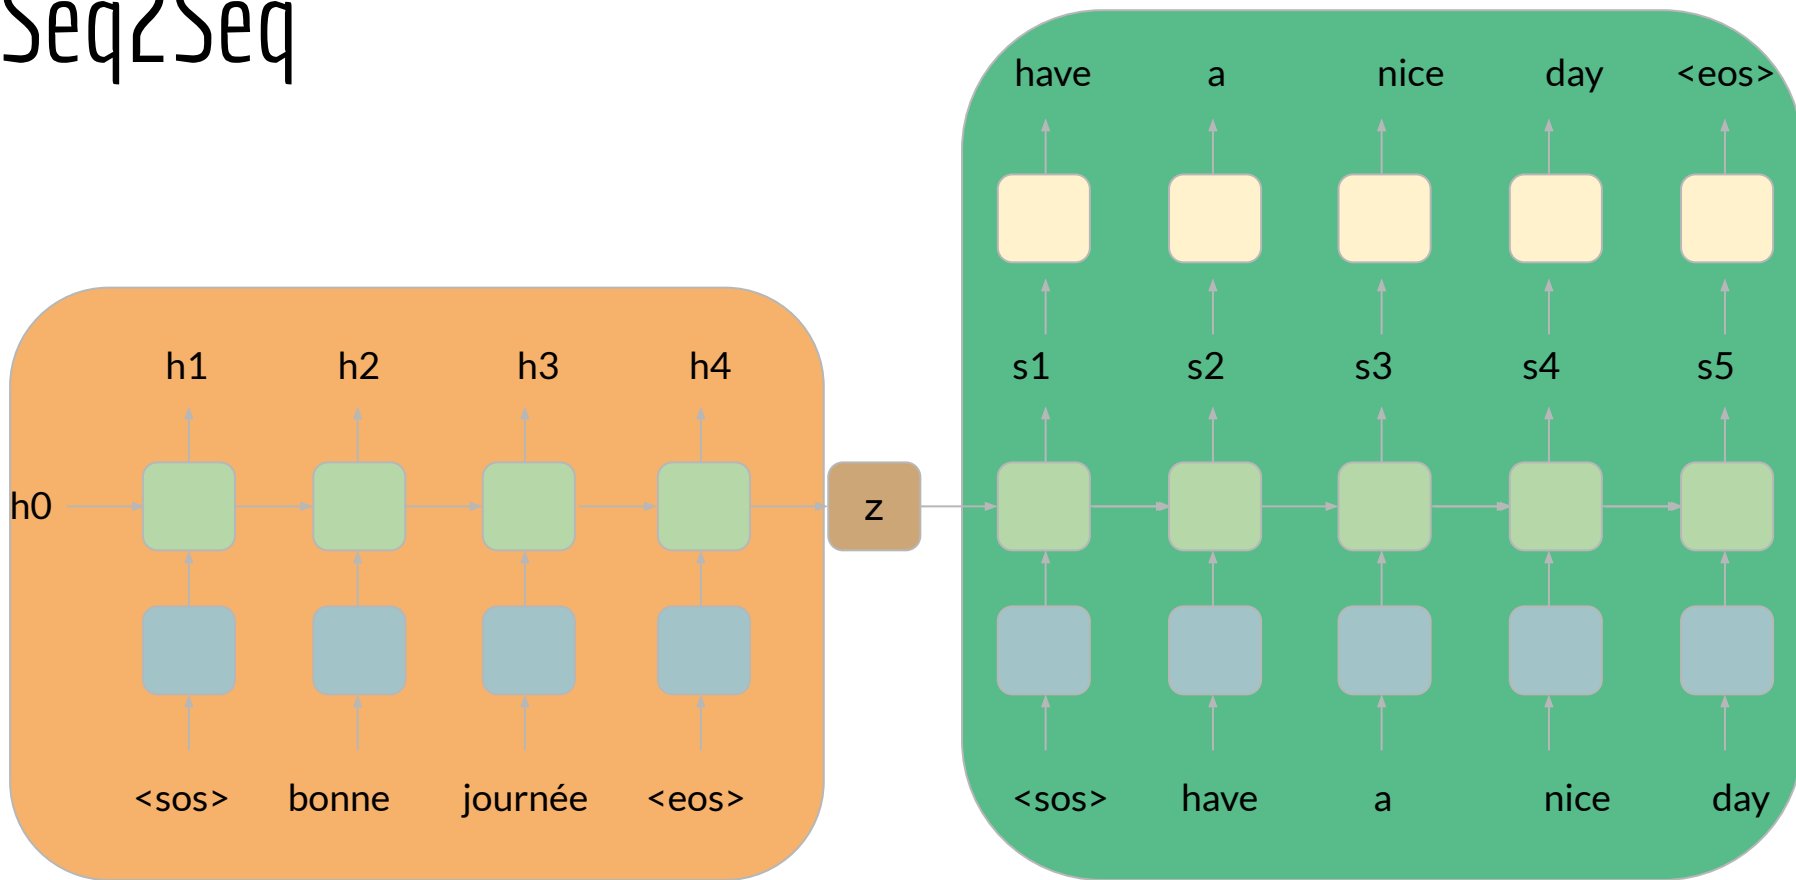
many to many



Seq2Seq

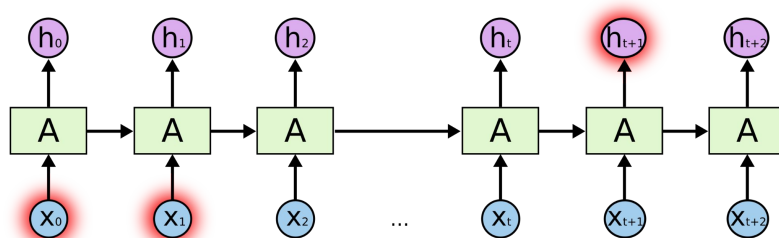


Seq2Seq



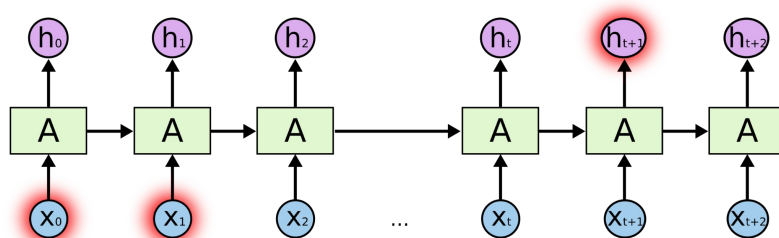
Sequence modeling

- Plain RNNs tend to perform poorly with very long sequences
- As information flows back through the network, it is lost or distorted

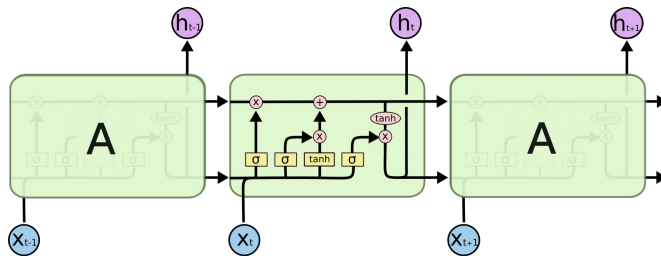


Sequence modeling

- Plain RNNs tend to perform poorly with very long sequences
- As information flows back through the network, it is lost or distorted

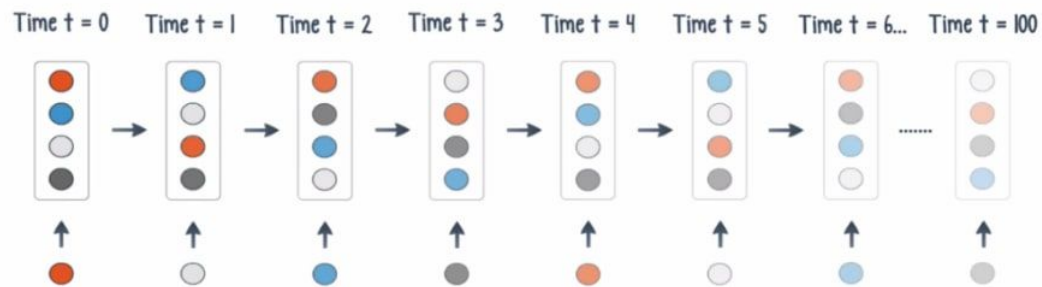


- LSTM cells introduce mechanisms that control the flow of information.



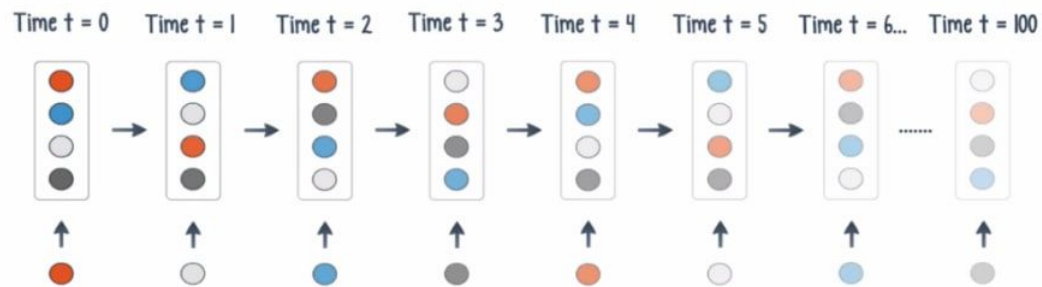
Problems with RNN

- Long Range Dependencies

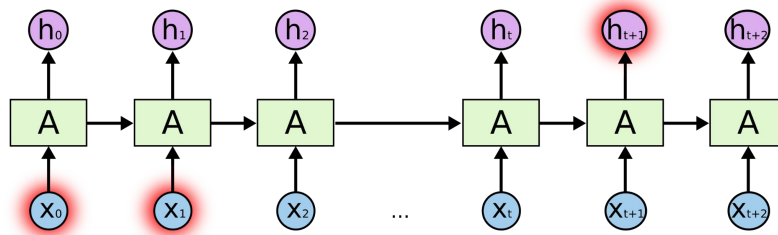


Problems with RNN

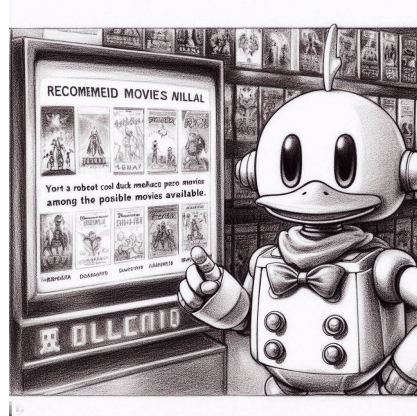
- Long Range Dependencies



- Parallelization



Introduction to Natural Language Processing



AI frameworks



Transformers



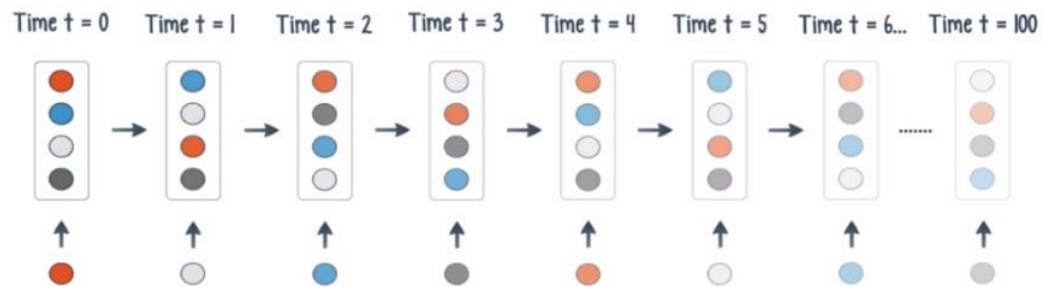
Disclaimer

Most of the illustrations and animations come from the great Jay Alammar blog post

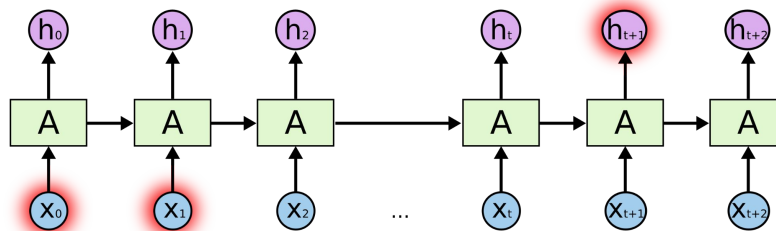
<http://jalammar.github.io/illustrated-transformer/>

Problems with RNN

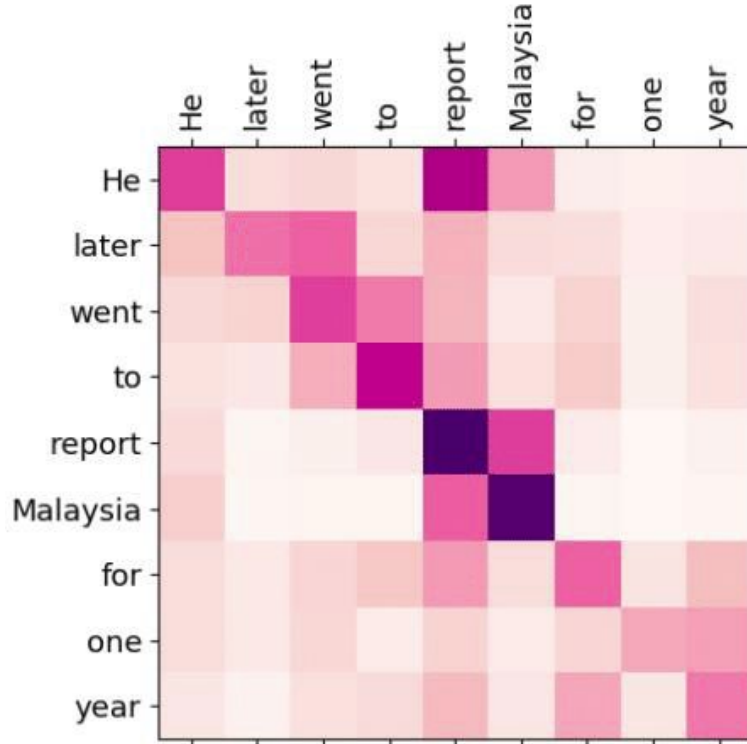
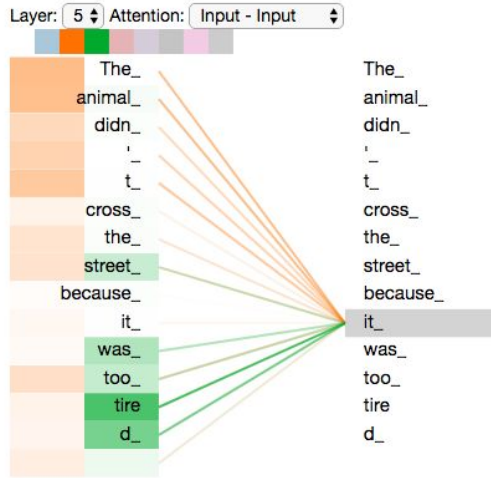
- Long Range Dependencies



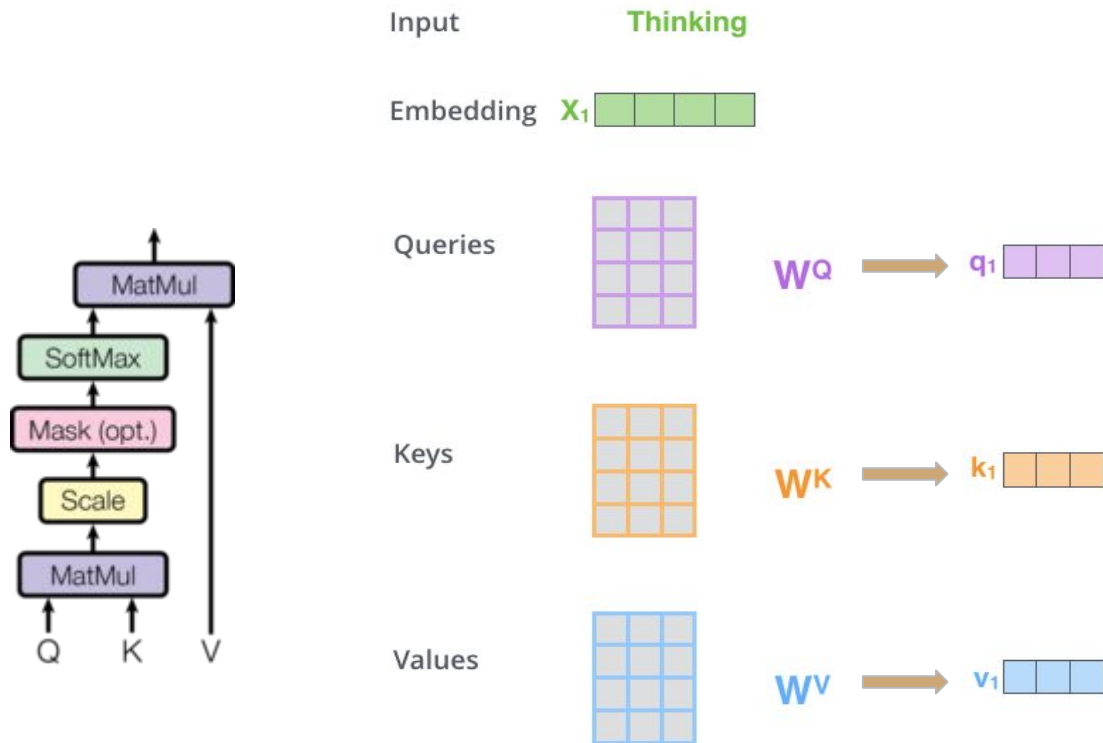
- Parallelization



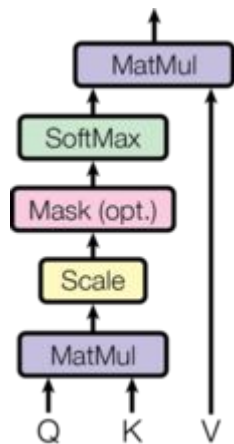
Solution: Attention and self attention

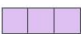


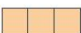
Self attention

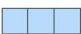


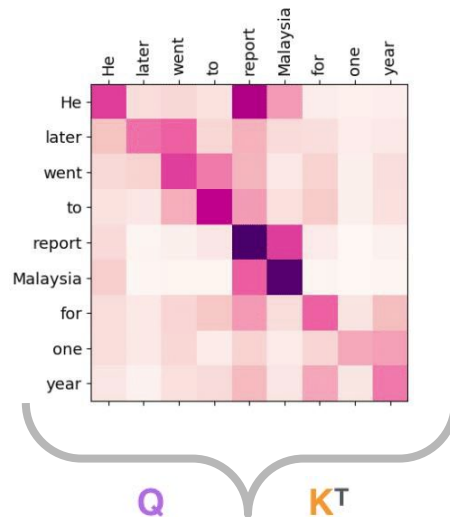
Self attention



q_1 

k_1 

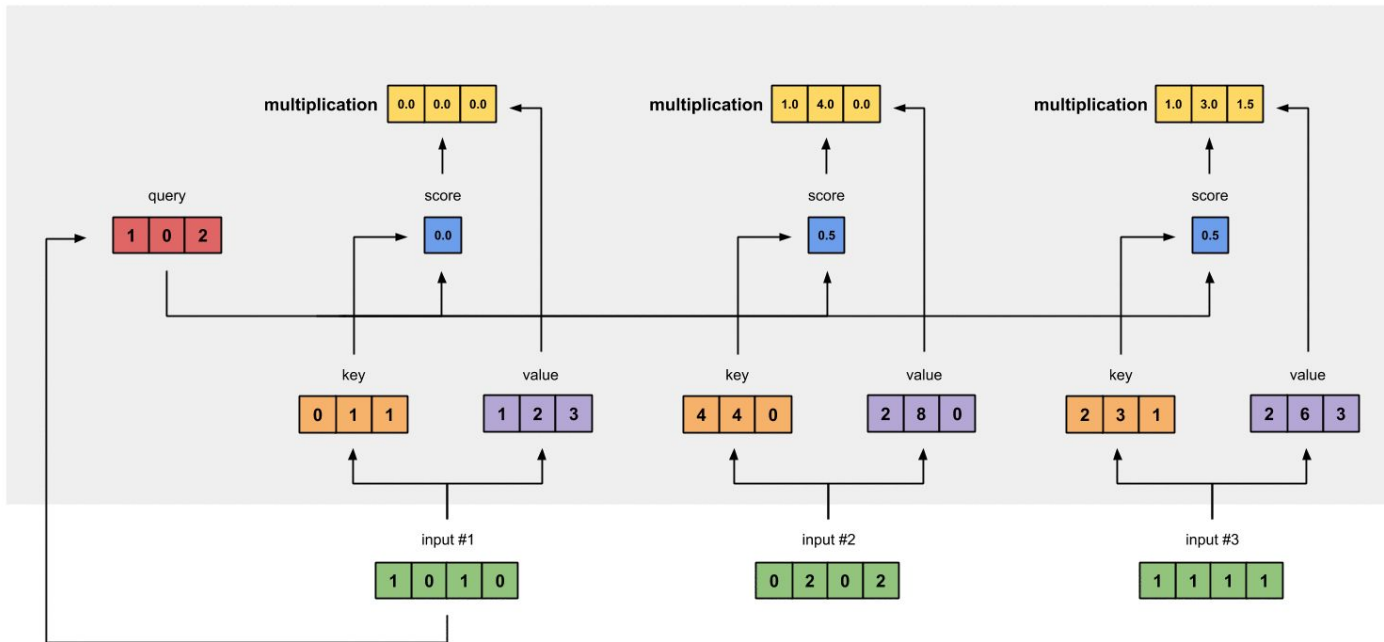
v_1 



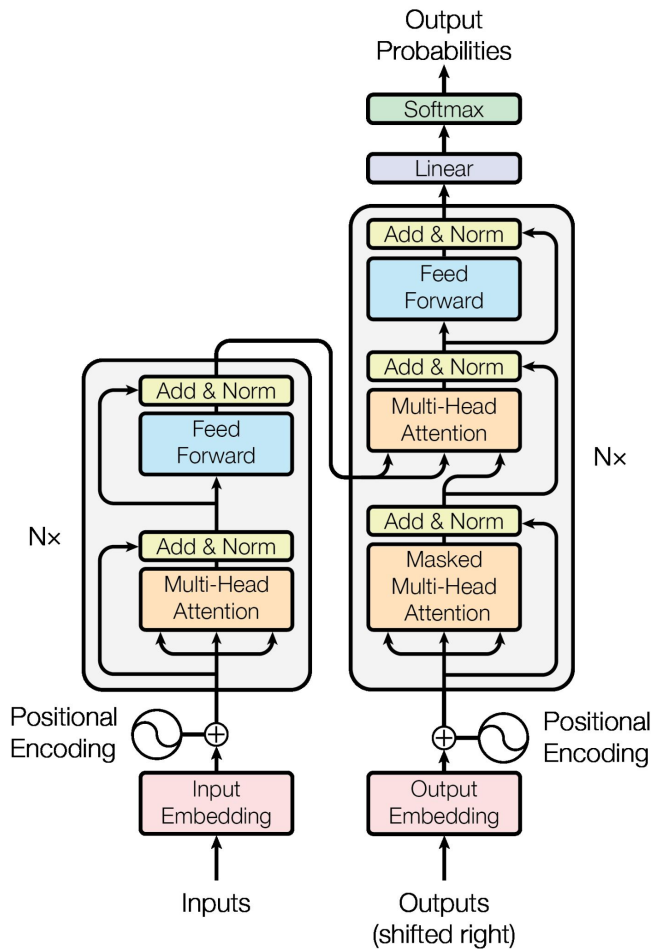
$$\text{softmax} \left(\frac{\begin{matrix} \text{3x3 purple grid} \\ \times \\ \text{3x3 orange grid} \end{matrix}}{\sqrt{d_k}} \right) \begin{matrix} \text{3x3 blue grid} \\ V \end{matrix}$$

$$= \begin{matrix} \text{3x3 pink grid} \\ Z \end{matrix}$$

Self attention

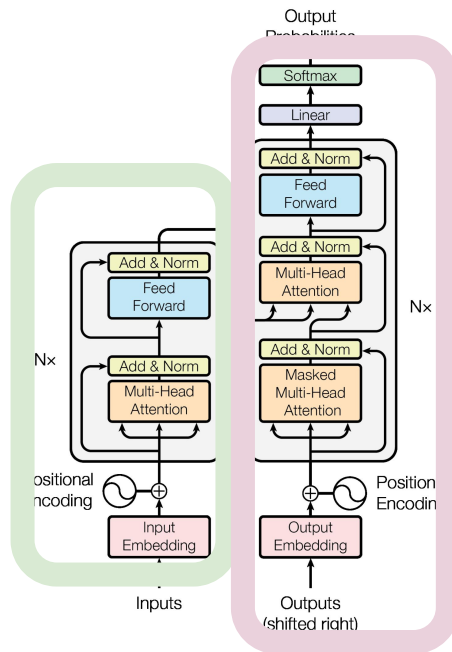
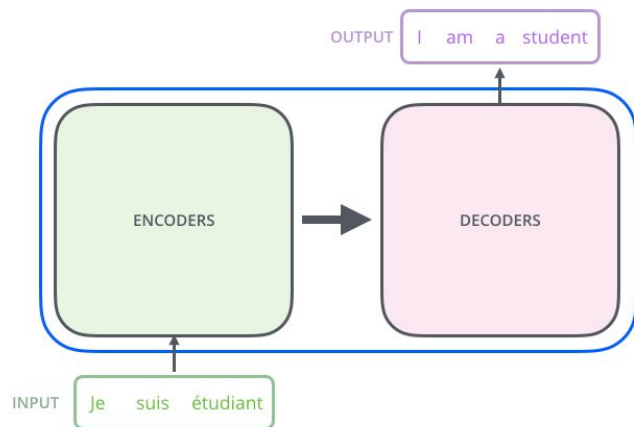


Transformers



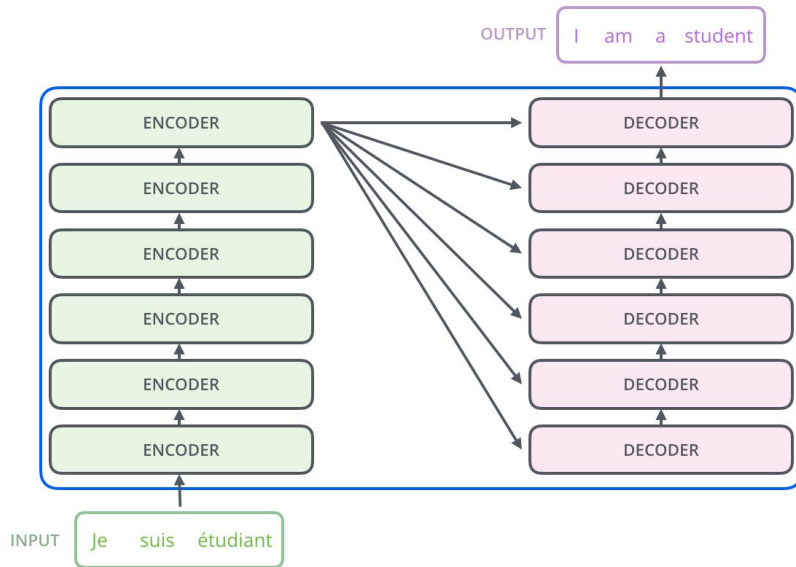
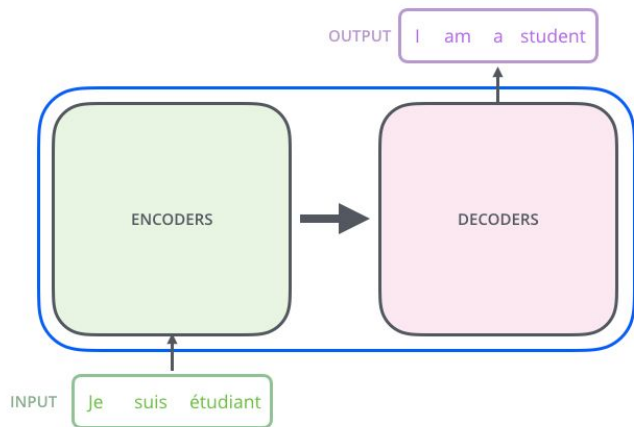
Transformers

- Encoder-decoder architecture

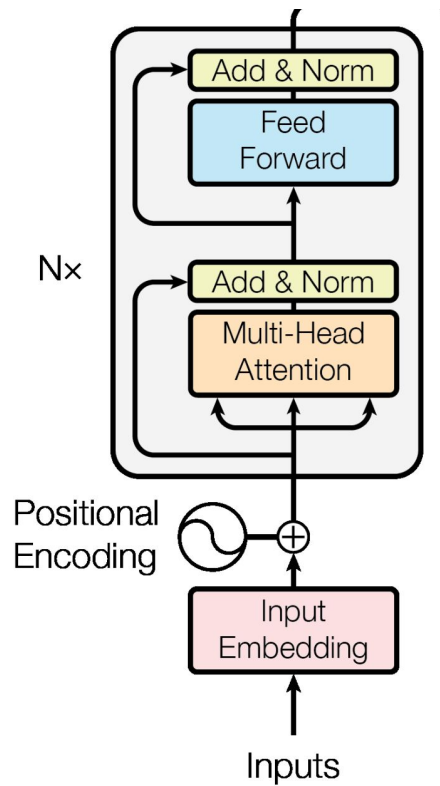


Transformers

- Encoder-decoder architecture

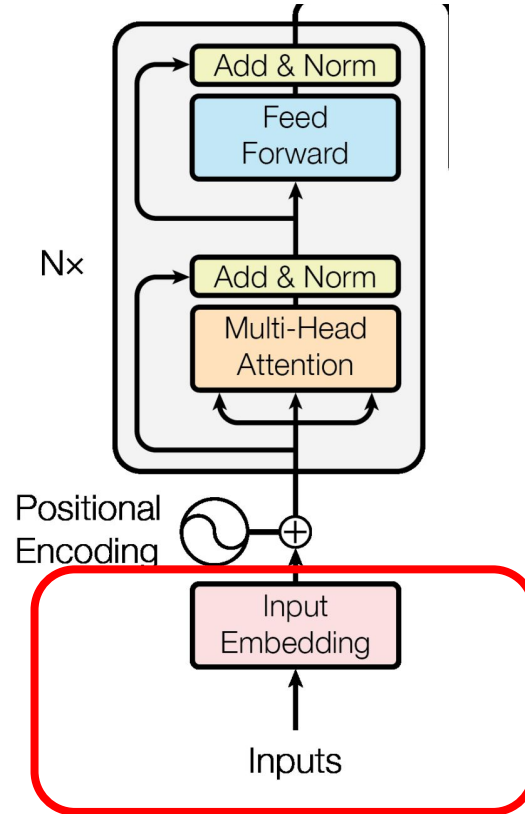


Encoder



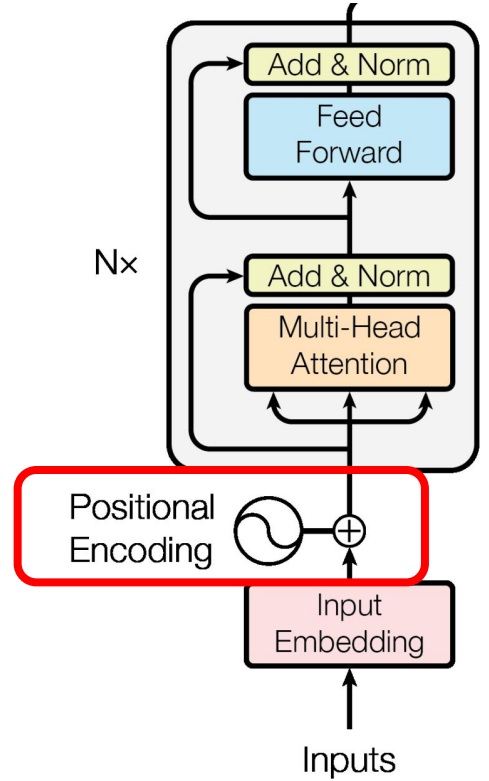
Encoder

- Input embeddings



Encoder

- Input embeddings
- Positional encoding



Encoder

- Input embeddings
- Positional encoding

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

Sequence Index of token Positional Encoding Matrix

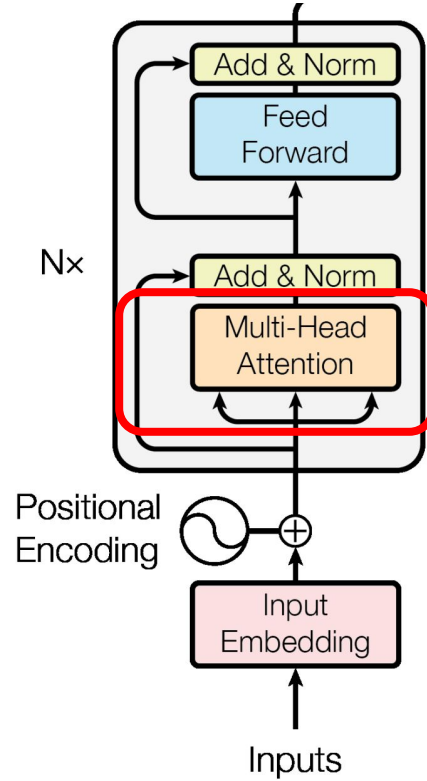
I	→	0	→	P ₀₀	P ₀₁	...	P _{0d}
am	→	1	→	P ₁₀	P ₁₁	...	P _{1d}
a	→	2	→	P ₂₀	P ₂₁	...	P _{2d}
Robot	→	3	→	P ₃₀	P ₃₁	...	P _{3d}

Positional Encoding Matrix for the sequence 'I am a robot'



Encoder

- Input embeddings
- Positional encoding
- Multi-head self attention



Encoder

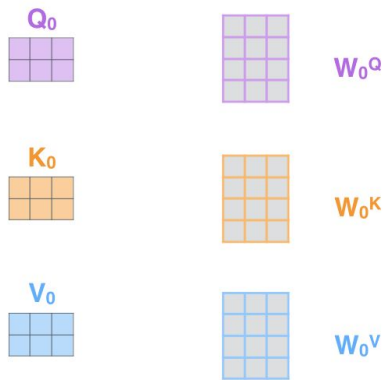
- Input embeddings
- Positional encoding
- Multi-head self attention

Thinking
Machines

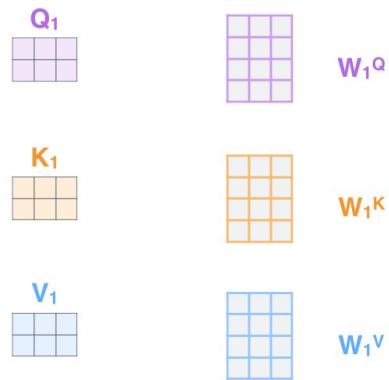
X



ATTENTION HEAD #0



ATTENTION HEAD #1



ATTENTION
HEAD #0



ATTENTION
HEAD #1



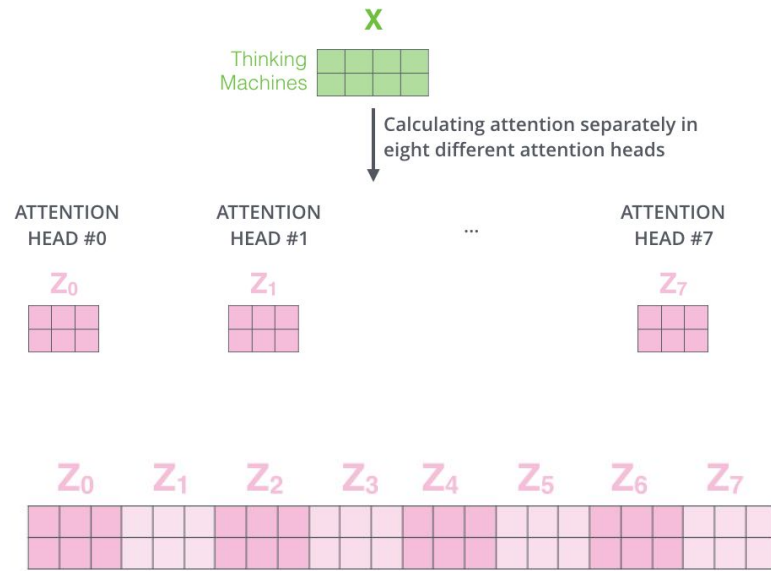
...

ATTENTION
HEAD #7



Encoder

- Input embeddings
- Positional encoding
- Multi-head self attention



Encoder

- Input embeddings
- Positional encoding
- Multi-head self attention

1) This is our input sentence*

Thinking
Machines

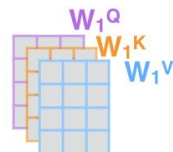
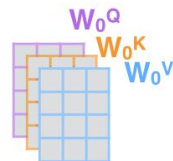
2) We embed each word*



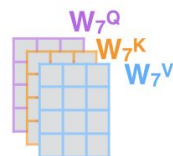
* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one



3) Split into 8 heads. We multiply X or R with weight matrices



...



4) Calculate attention using the resulting $Q/K/V$ matrices



...



5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer



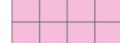
...



W^O

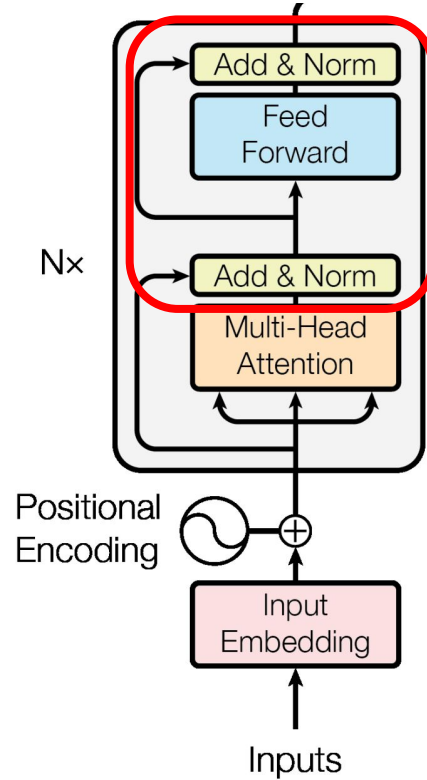


Z



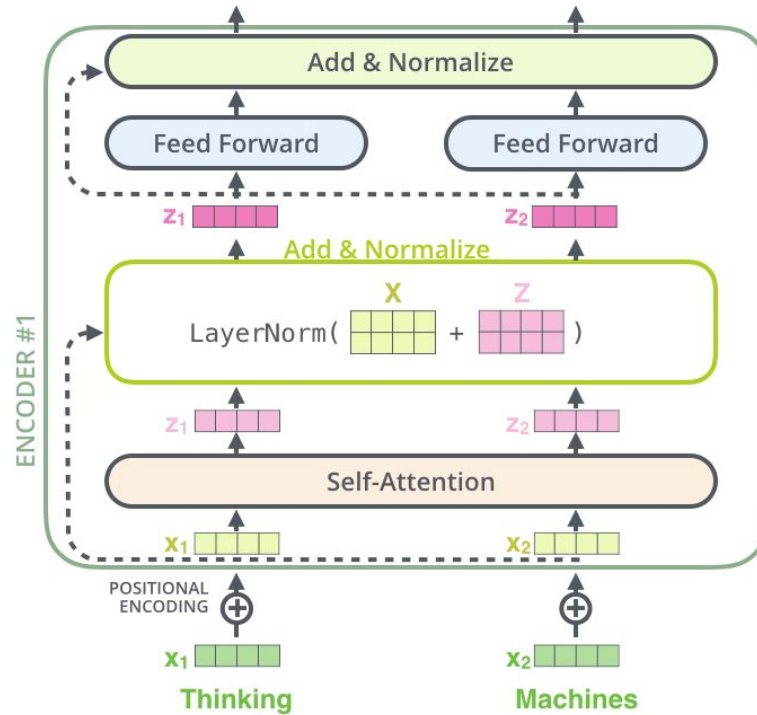
Encoder

- Input embeddings
- Positional encoding
- Multi-head self attention
- Residual connections and normalization
- Feed Forward neural net
- Residual connections and normalization



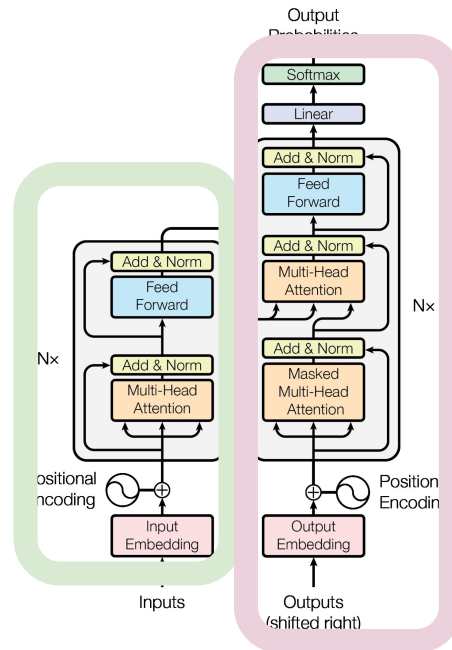
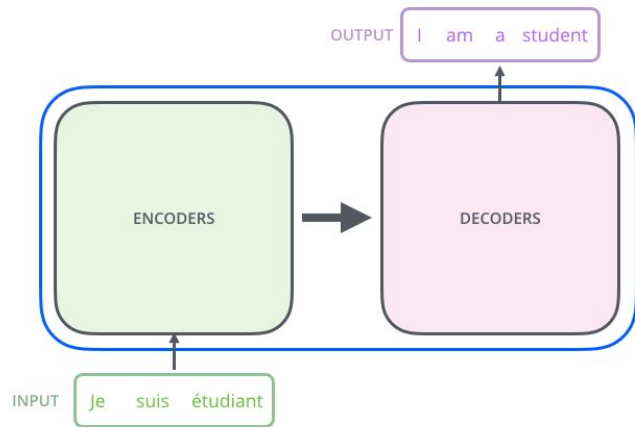
Encoder

- Input embeddings
- Positional encoding
- Multi-head self attention
- Residual connections and normalization
- Feed Forward neural net
- Residual connections and normalization

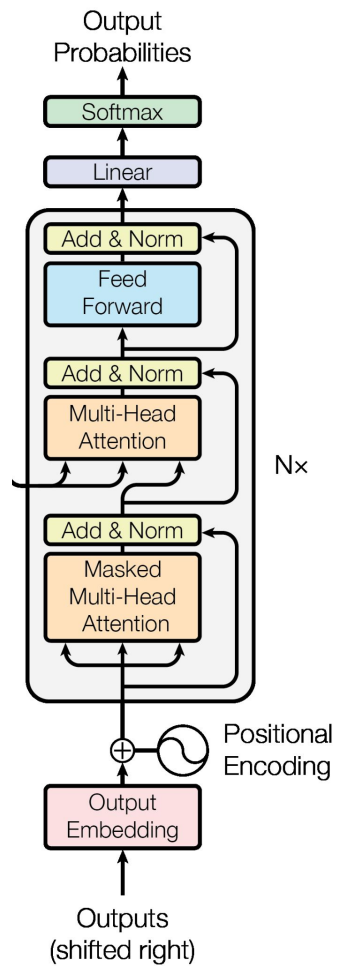


Transformers

- Encoder-decoder architecture

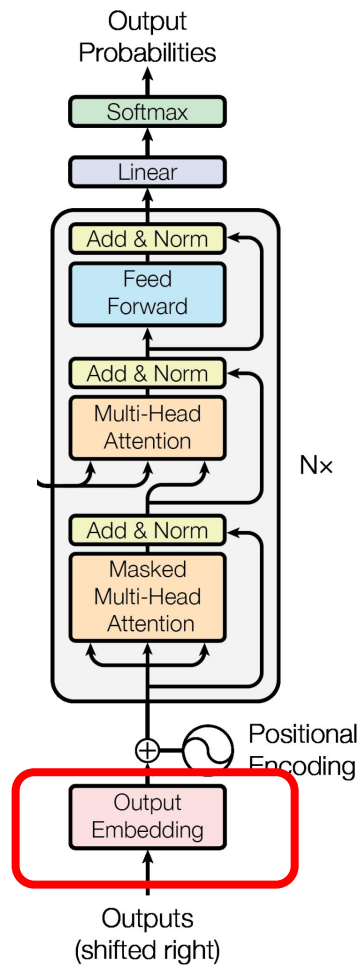


Decoder



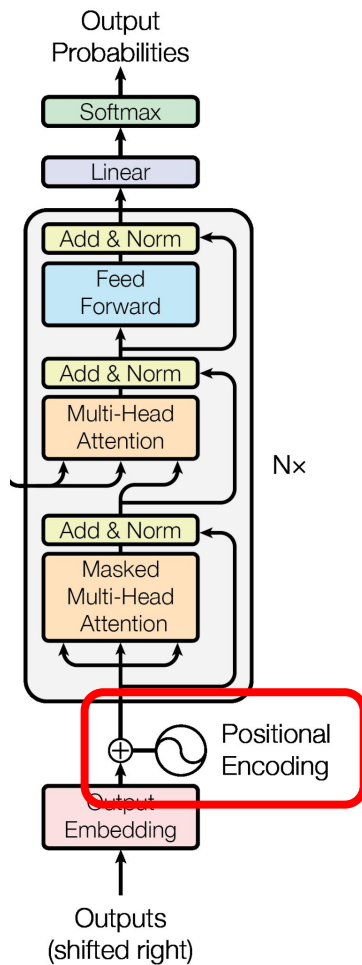
Decoder

- Output embeddings



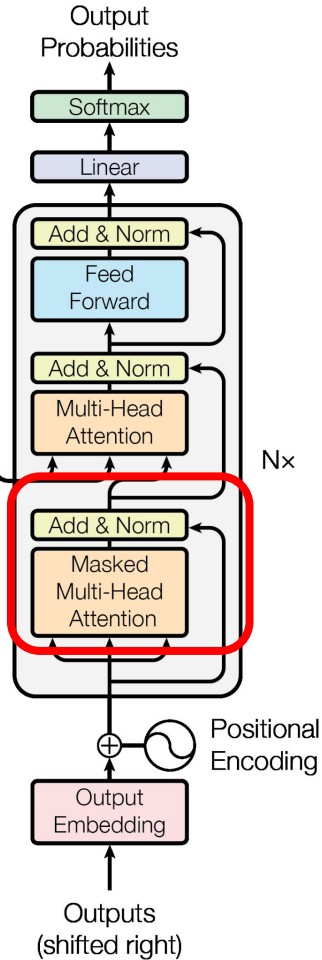
Decoder

- Output embeddings
- Positional encoding



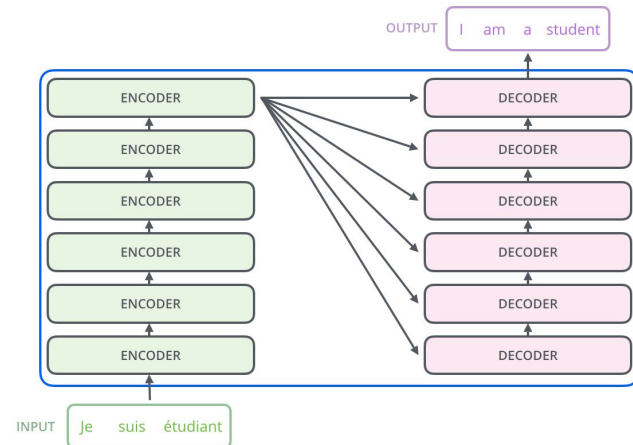
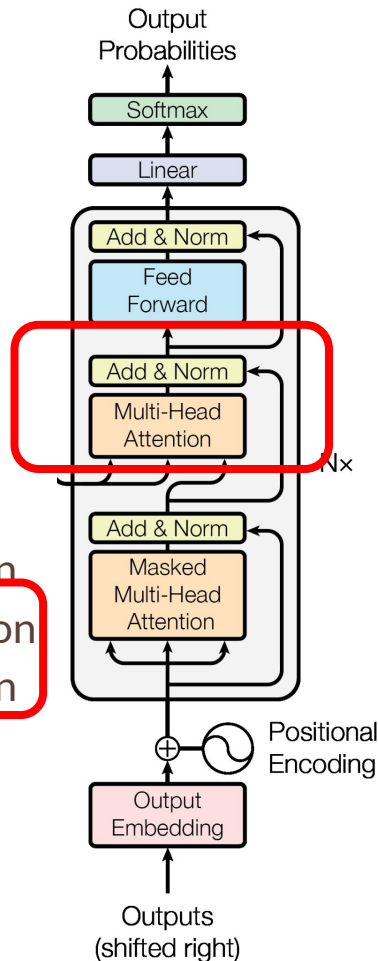
Decoder

- Output embeddings
- Positional encoding
- Masked Multi-head self attention
- Residual connections & normalization



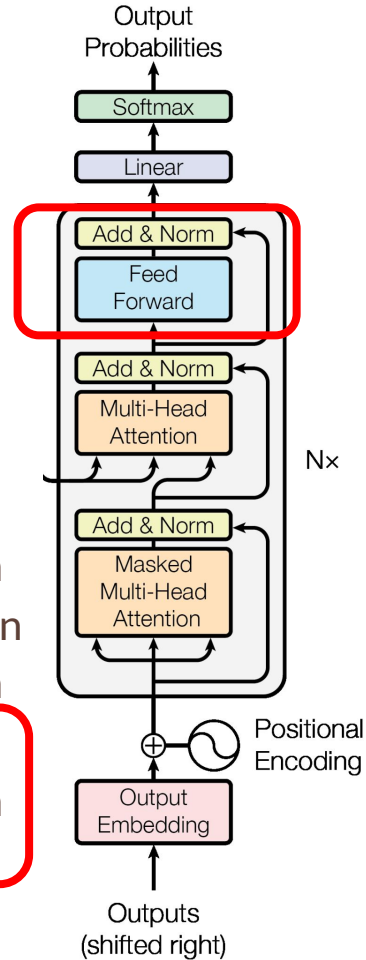
Decoder

- Output embeddings
- Positional encoding
- Masked Multi-head self attention
- Residual connections & normalization
- Multi-head encoder-decoder attention
- Residual connections & normalization



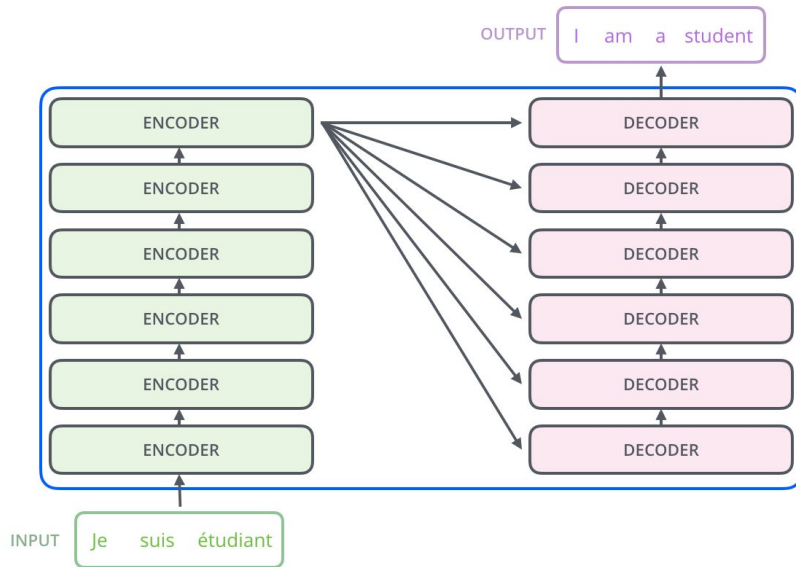
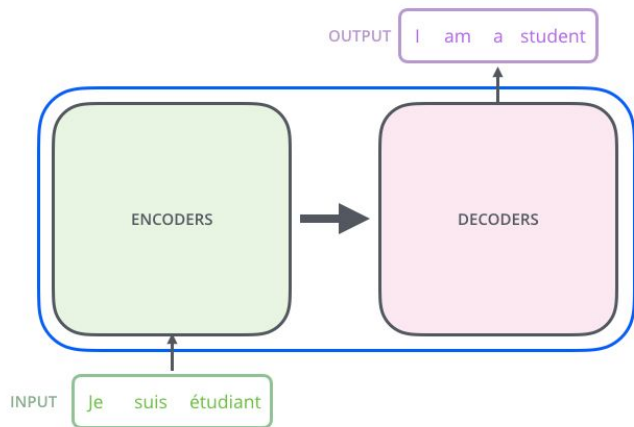
Decoder

- Output embeddings
- Positional encoding
- Masked Multi-head self attention
- Residual connections & normalization
- Multi-head encoder-decoder attention
- Residual connections & normalization
- Feed forward neural network
- Residual connections & normalization

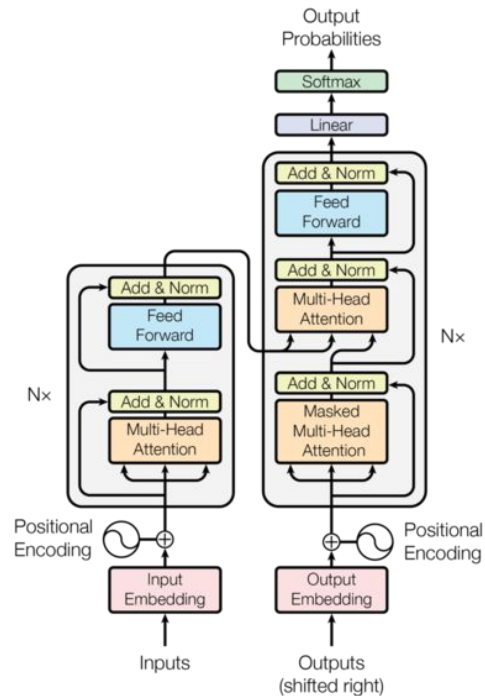


Transformers

- Encoder-decoder architecture

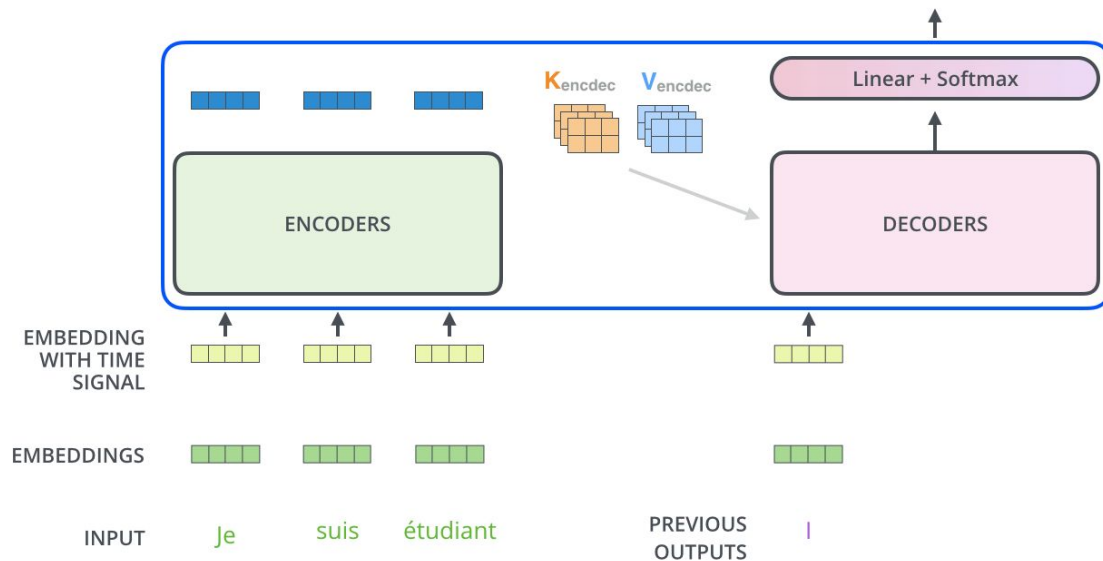


Transformers



Decoding time step: 1 2 3 4 5 6

OUTPUT |

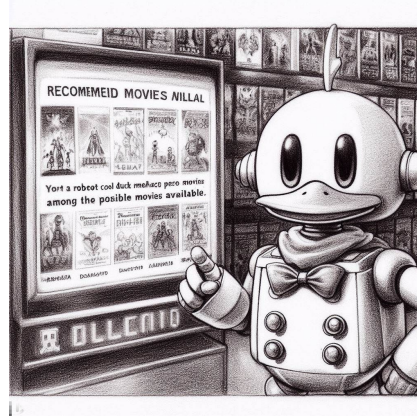


Go to this blog

Jay Alammar blog post: The illustrated transformer

<http://jalammar.github.io/illustrated-transformer/>

Introduction to Natural Language Processing



AI frameworks



Large Language models



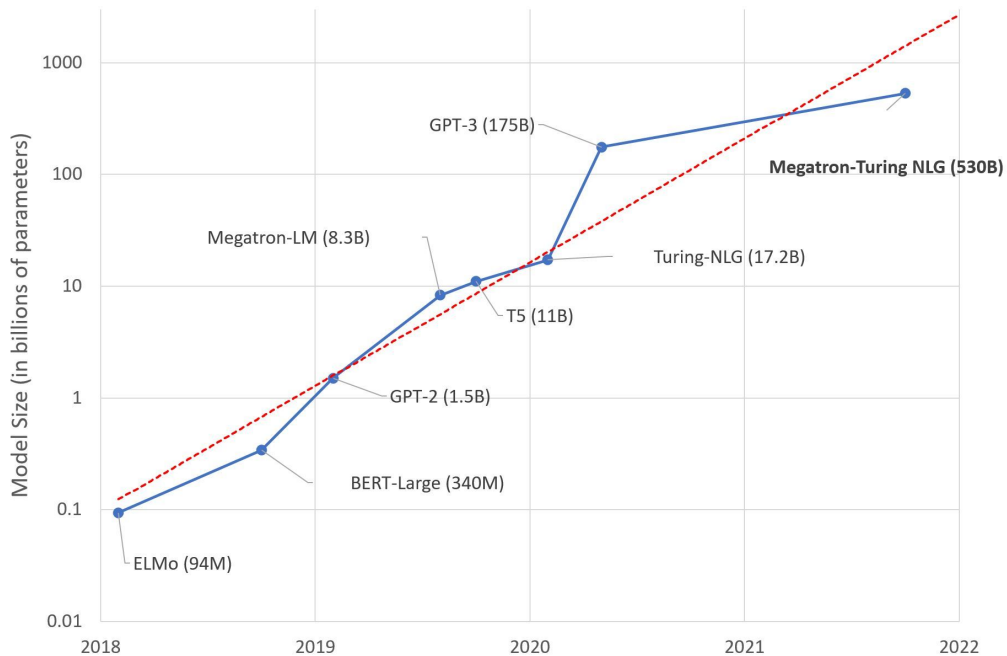
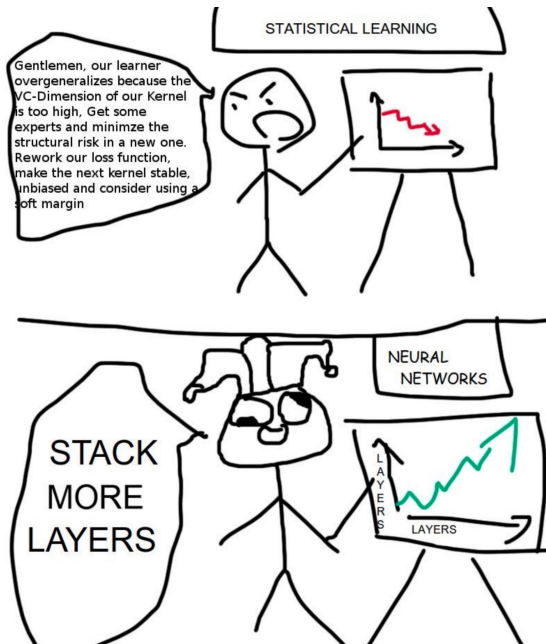
Disclaimer

All of these slides are directly modified from ones created by Adil Zouitine

Large Language Model

- **BERT:**
 - Masked Language Modeling (MLM)
 - Next Sentence Prediction (NSP)
 - Needs fine tuning!
- **GPT :**
 - Language modeling: predict the next word
- **T5 (Text-to-Text Transfer Transformer):**
 - denoising auto-encoding task

Large Language Model



Generative Pretrain Transformer aka GPT

GPT (2018) : Improving Language Understanding by Generative Pre-Training

TLDR: Generative LLM + Academic dataset

GPT 2 (2019): Language Models are Unsupervised Multitask Learners

TLDR: 100 x Bigger than GPT + Generative LLM + Common crawl dataset

GPT 3 (2020): Language Models are Few-Shot Learners

TLDR: 100 x Bigger than GPT 2 + Generative LLM + Common crawl dataset

GPT 4 (2023) :

TLDR: 400 x Bigger than GPT 3 + Generative LLM + Common crawl (img + txt)



Chat GPT is a cake

How Much Information is the Machine Given during Learning?

Y. LeCun

► **“Pure” Reinforcement Learning (cherry)**

- The machine predicts a scalar reward given once in a while.

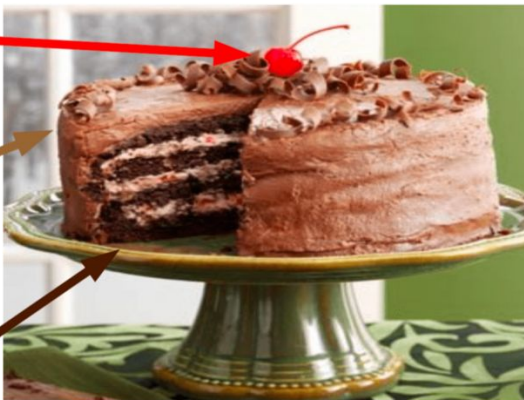
► **A few bits for some samples**

► **Supervised Learning (icing)**

- The machine predicts a category or a few numbers for each input
- Predicting human-supplied data
- **10→10,000 bits per sample**

► **Self-Supervised Learning (cake génoise)**

- The machine predicts any part of its input for any observed part.
- Predicts future frames in videos
- **Millions of bits per sample**



Reinforcement learning from human feedback

Step 1

Collect demonstration data and train a supervised policy.

A prompt is sampled from our prompt dataset.

A labeler demonstrates the desired output behavior.

This data is used to fine-tune GPT-3.5 with supervised learning.



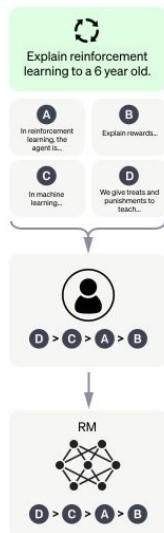
Step 2

Collect comparison data and train a reward model.

A prompt and several model outputs are sampled.

A labeler ranks the outputs from best to worst.

This data is used to train our reward model.



Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.

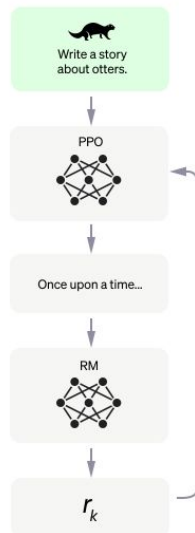
A new prompt is sampled from the dataset.

The PPO model is initialized from the supervised policy.

The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.



Future of LLM

- Small open source foundation models: LLaMA2 , Mistral, ...
- Fit on single computers
- Fine-tune for specific usages
- Comparable performance than ChatGPT 3.5