

Sequence Modeling: Recurrent Neural Networks

High Dimensional and Deep Learning – juliette.chevallier@insa-toulouse.fr
N7 & INSA Toulouse, 5th year ModIA

1. Sequence Modeling
2. Recurrent Neural Networks
3. Training Recurrent Networks
4. Challenge of Long-Term Dependencies
5. Gated Recurrent Neural Network
6. Explicit Memory
7. Appendix: Categorical Variable Encoding

Sequence Modeling


1.1 Sequential Data

1.2 Traditional Time Series Models

Sequential Data


Sequence ↔ Explicit **order** on the observations that must be preserved when training models and making predictions.

- **Sequence Prediction:** Weather forecasting, Stock market prediction, Product recommendation
- **Sequence Classification:** DNA seq. classification, Anomaly detection, Sentiment analysis;
- **Sequence Generation:** Text generation, Handwriting prediction, Music generation;
- **Sequence-to-Sequence Prediction:** Multi-Step time series forecasting, Text summarization, Program execution.



From: Vinyals, Toshev, Bengio, Erhan. Show and tell: A neural image caption generator. CVPR 2015

Do not get rid of traditional models too quickly!




Baughman, Haas, Wolski, Foster, Chard. *Predicting Amazon Spot Prices with LSTM Networks*. 2018.

Sequence Modeling

1.1 Sequential Data

1.2 Traditional Time Series Models

Time Series Forecasting



AR Autoregressive

MA Moving average

ARMA Autoregressive moving average

ARIMA Autoregressive integrated moving average

SARIMA Seasonal autoregressive integrated moving average

- Describing temporal dynamics in great detail;
- Realization of a stochastic process;
- Decomposition: trend, seasonality and (stochastic) remainder.

ARIMA vs. (Recurrent) Neural Network

ARIMA:

- + Out-perform deep learning methods for forecasting **short-term**
- + Out-perform deep learning methods for forecasting on **univariate** data

Neural Network:

ARIMA vs. (Recurrent) Neural Network

ARIMA:

- + Out-perform deep learning methods for forecasting **short-term**
- + Out-perform deep learning methods for forecasting on **univariate** data

Neural Network:

- Focus on univariate data, with linear relationships, and fixed and manually-diagnosed temporal dependence

ARIMA vs. (Recurrent) Neural Network

ARIMA:

- + Out-perform deep learning methods for forecasting **short-term**
- + Out-perform deep learning methods for forecasting on **univariate** data
- Focus on univariate data, with linear relationships, and fixed and manually-diagnosed temporal dependence

Neural Network:

- + Ability to learn noisy and non-linear relationships, regardless of the number of inputs
- + Relax univariate assumption

ARIMA vs. (Recurrent) Neural Network

ARIMA:

- + Out-perform deep learning methods for forecasting **short-term**
- + Out-perform deep learning methods for forecasting on **univariate** data
- Focus on univariate data, with linear relationships, and fixed and manually-diagnosed temporal dependence

Neural Network:

- + Ability to learn noisy and non-linear relationships, regardless of the number of inputs
- + Relax univariate assumption
- More complicated and difficult to train
- Do not exceed the performance of ARIMA in most cases

ARIMA vs. (Recurrent) Neural Network

ARIMA:

- + Out-perform deep learning methods for forecasting **short-term**
- + Out-perform deep learning methods for forecasting on **univariate** data
- Focus on univariate data, with linear relationships, and fixed and manually-diagnosed temporal dependence

Neural Network:

- + Ability to learn noisy and non-linear relationships, regardless of the number of inputs
- + Relax univariate assumption
- More complicated and difficult to train
- Do not exceed the performance of ARIMA in most cases

~> Neural Networks

Sequential Data: Audio, Speech, Language, Videos with time context, Time series

Question: *How can we integrate this context in neural networks?*

Naive approaches

Sequential Data: Audio, Speech, Language, Videos with time context, Time series

***Question:** How can we integrate this context in neural networks?*

Multilayer Perceptron Regression:

Time series prediction \longleftrightarrow Regression problem: x_{t+1} as a function of x_t

Sequential Data: Audio, Speech, Language, Videos with time context, Time series

***Question:** How can we integrate this context in neural networks?*

Multilayer Perceptron Regression:

Time series prediction \longleftrightarrow **Regression** problem: x_{t+1} as a function of x_t
 \leadsto *Multilayer Perceptron model*

- Difficult **training**,
- No more efficient than an ARIMA model (or even less)

Sequential Data: Audio, Speech, Language, Videos with time context, Time series

***Question:** How can we integrate this context in neural networks?*

Multilayer Perceptron Regression:

Time series prediction \longleftrightarrow **Regression** problem: x_{t+1} as a function of x_t
 \leadsto *Multilayer Perceptron model*

- Difficult **training**,
- No more efficient than an ARIMA model (or even less)

Huge Convolutional Network:

Feed the whole sequence to a huge network

- Inefficient memory usage,
- Difficult/Impossible to **train**,
- Difference between spatial and temporal dimensions?
- Not real-time (translation !)

Naive approaches

Sequential Data: Audio, Speech, Language, Videos with time context, Time series

***Question:** How can we integrate this context in neural networks?*

Multilayer Perceptron Regression:

Time series prediction \longleftrightarrow Regression problem: x_{t+1} as a function of x_t
 \leadsto Multilayer Perceptron model

- Difficult training,
- No more efficient than an ARIMA model (or even less)

Huge Convolutional Network:

Feed the whole sequence to a huge network

- Inefficient memory usage
- Difficult/Impossible to train,
- Difference between spatial and temporal dimensions?
- Not real-time (translation !)

Do not do that!

Naive approaches

Sequential Data: Audio, Speech, Language, Videos with time context, Time series

***Question:** How can we integrate this context in neural networks?*

Multilayer Perceptron Regression:

Time series prediction \longleftrightarrow **Regression** problem: x_{t+1} as a function of x_t
 \leadsto *Multilayer Perceptron model*

- Difficult **training**,
- No more efficient than an ARIMA model (or even less)

Huge Convolutional Network:

Feed the whole sequence to a huge network

- Inefficient memory usage
- Difficult/Impossible to **train**,
- Difference between spatial and temporal dimensions?
- Not real-time (translation !)

Do not do that!

\leadsto **Recurrent Neural Networks**

1. Sequence Modeling

- 1.1 Sequential Data
- 1.2 Traditional Time Series Models

2. Recurrent Neural Networks

- 2.1 Recurrent Neural Networks
- 2.2 Bidirectional RNN
- 2.3 Encoder-Decoder Seq-2-Seq Architectures

3. Training Recurrent Networks

- 3.1 Forward Propagation
- 3.2 Back-Propagation Through Time
- 3.3 Limit of the simple RNN model

4. Challenge of Long-Term Dependencies

- 4.1 Optimization for Long-Term Dependencies
- 4.2 The Challenge of Long-Term Dependencies

5. Gated Recurrent Neural Network

- 5.1 Long-Short-Term-Memory Architecture
- 5.2 Variants on Long-Short-Term-Memory
- 5.3 Gated Recurrent Unit

6. Explicit Memory

- 6.1 Memory Networks
- 6.2 Neural Turing Machines

7. Appendix: Categorical Variable Encoding

Recurrent Neural Networks

2.1 Recurrent Neural Networks

2.2 Bidirectional RNN

2.3 Encoder-Decoder Seq-2-Seq Architectures

Information Persistence – Parameter-Sharing

- Humans don't start their thinking from scratch every second:
While reading, each word is understood according to your previous ones.
Your thoughts have persistence.
- Traditional neural networks **can't** do this.

Information Persistence – Parameter-Sharing

- Humans don't start their thinking from scratch every second:
While reading, each word is understood according to your previous ones.
Your thoughts have persistence.

- Traditional neural networks **can't** do this.

~> **Recurrent Neural Networks:**


Networks with loops in them, allowing **information to persist**.

Information Persistence – Parameter-Sharing

- Humans don't start their thinking from scratch every second:
While reading, each word is understood according to your previous ones.
Your thoughts have persistence.
- Traditional neural networks **can't** do this.

↪ Recurrent Neural Networks:

Networks with loops in them, allowing **information to persist**.



Input x_t at time t


Output h_t at time t

Loop allows information to be passed from one step of the network to the next

*Little (1974), Hopfield (1982),
Rumelhart, Hinton & Williams (1986), Elman (1990)*

Unfolding Computational Graphs


- RNN \equiv Multiple copies of the **same** network,
each passing a message to its successor: $h_t = f(h_{t-1}, x_t; \theta)$.



- **Advantages of the unfolding process:**
 - Whatever the sequence length, the learned model always has the same input size,
 - Possible to use the **same** transition f with same parameters at every time step.

Unfolding Computational Graphs

- RNN \equiv Multiple copies of the **same** network,
each passing a message to its successor: $h_t = f(h_{t-1}, x_t; \theta)$.




- **Advantages of the unfolding process:**

- Whatever the sequence length, the learned model always has the same input size,
- Possible to use the **same** transition f with same parameters at every time step.

\leadsto A **single** model f that operates

- on **all** time steps,
- and **all** sequence lengths.



Standard Recurrent Neural Networks

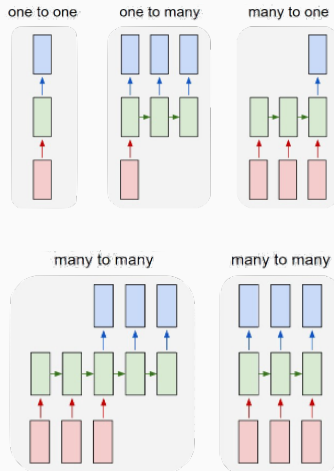
Different types of RNN's:

- One-to-one e.g. Image classification,
- One-to-Many e.g. Image captioning,
- Many-to-One e.g. Sentiment analysis,
- Many-to-Many e.g. Machine translation.

The brain worm polka



themachinefolksession.org



RNNs are Turing-Complete

RNN:

1. combine the input vector
2. with their state vector
3. using a fixed (but learned) function
4. to produce a new state vector

RNNs are Turing-Complete

RNN:

1. combine the input vector
2. with their state vector
3. using a fixed (but learned) function
4. to produce a new state vector

Programming terms: Running a fixed program with certain inputs and some internal variables

RNNs essentially describe programs

RNNs are Turing-Complete

RNN:

1. combine the input vector
2. with their state vector
3. using a fixed (but learned) function
4. to produce a new state vector

RNNs are Turing-Complete

They can simulate arbitrary programs
(with proper weights)

Programming terms: Running a fixed program with certain inputs and some internal variables

RNNs essentially describe programs

- Similar to universal approximation theorems for neural nets
- Shouldn't read too much into this.

Siegelmann and Sontag (1992)

Recurrent Neural Networks

2.1 Recurrent Neural Networks

2.2 Bidirectional RNN

2.3 Encoder-Decoder Seq-2-Seq Architectures

Bidirectional RNNs

- All the recurrent networks we have considered so far have a **causal** structure:
State at time t captures only information from the past x_1, \dots, x_{t-1} , and the present input x_t
- In many *applications*, however, we want to output a prediction of y_t that may depend on the **whole** input sequence
Handwriting recognition, Speech recognition, Bio-informatics, etc.

Bidirectional RNNs


- All the recurrent networks we have considered so far have a **causal** structure:
State at time t captures only information from the past x_1, \dots, x_{t-1} , and the present input x_t
- In many *applications*, however, we want to output a prediction of y_t that may depend on the **whole** input sequence
Handwriting recognition, Speech recognition, Bio-informatics, etc.

↪ **Bidirectional** recurrent neural networks

Bidirectional RNNs

- All the recurrent networks we have considered so far have a **causal** structure:
State at time t captures only information from the past x_1, \dots, x_{t-1} , and the present input x_t
- In many *applications*, however, we want to output a prediction of y_t that may depend on the **whole** input sequence
Handwriting recognition, Speech recognition, Bio-informatics, etc.

↪ **Bidirectional** recurrent neural networks




Schuster & Paliwal (1997), Graves (2012)

Bidirectional RNNs

- All the recurrent networks we have considered so far have a **causal** structure:
State at time t captures only information from the past x_1, \dots, x_{t-1} , and the present input x_t
- In many *applications*, however, we want to output a prediction of y_t that may depend on the **whole** input sequence

Handwriting recognition, Speech recognition, Bio-informatics, etc.

↪ **Bidirectional** recurrent neural networks




Schuster & Paliwal (1997), Graves (2012)

Bidirectional RNNs

- All the recurrent networks we have considered so far have a **causal** structure:
State at time t captures only information from the past x_1, \dots, x_{t-1} , and the present input x_t
- In many *applications*, however, we want to output a prediction of y_t that may depend on the **whole** input sequence

Handwriting recognition, Speech recognition, Bio-informatics, etc.

↪ Bidirectional recurrent neural networks




Schuster & Paliwal (1997), Graves (2012)

Bidirectional RNNs

- All the recurrent networks we have considered so far have a **causal** structure:
State at time t captures only information from the past x_1, \dots, x_{t-1} , and the present input x_t
- In many *applications*, however, we want to output a prediction of y_t that may depend on the **whole** input sequence

Handwriting recognition, Speech recognition, Bio-informatics, etc.

↪ Bidirectional recurrent neural networks



Schuster & Paliwal (1997), Graves (2012)

Recurrent Neural Networks


2.1 Recurrent Neural Networks

2.2 Bidirectional RNN

2.3 Encoder-Decoder Seq-2-Seq Architectures

Sequence-to-Sequence Model

The model consists of **3 parts**: **Encoder**, **Context** (encoder vector) and **Decoder**.




*Cho, Van Merriënboer, Gulcehre, Bahdanau, Bougares, Schwenk & Bengio (2014),
Sutskever, Vinyals & Le (2014)*

Sequence-to-Sequence Model

The model consists of **3 parts**: **Encoder**, **Context** (encoder vector) and **Decoder**.

Encoder:

- Stack of *several* recurrent units,
- Each accepts a *single* element of the input sequence,
- Propagates information *forward*: $h_t = f(h_{t-1}, x_t; \theta)$.



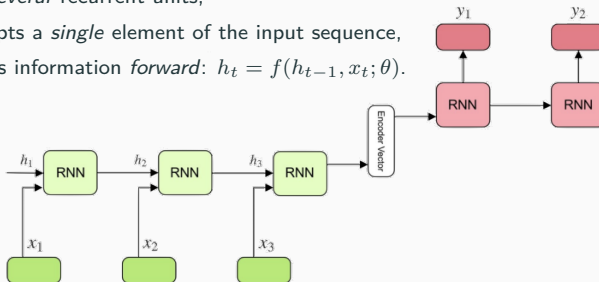
Cho, Van Merriënboer, Gulcehre, Bahdanau, Bougares, Schwenk & Bengio (2014),
Sutskever, Vinyals & Le (2014)

Sequence-to-Sequence Model

The model consists of **3 parts**: **Encoder**, **Context** (encoder vector) and **Decoder**.

Encoder:

- Stack of *several* recurrent units,
- Each accepts a *single* element of the input sequence,
- Propagates information *forward*: $h_t = f(h_{t-1}, x_t; \theta)$.



Context:

- Final hidden state produced from the encoder
- Goal: Encapsulate the information for all input elements

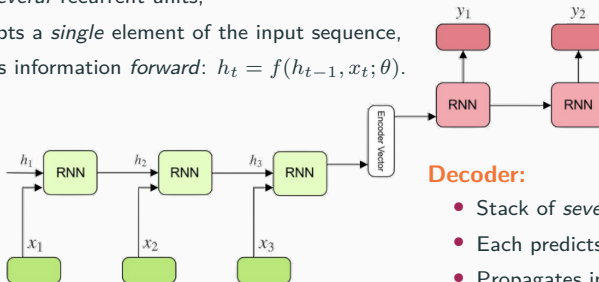
*Cho, Van Merriënboer, Gulcehre, Bahdanau, Bougares, Schwenk & Bengio (2014),
Sutskever, Vinyals & Le (2014)*

Sequence-to-Sequence Model

The model consists of **3 parts**: **Encoder**, **Context** (encoder vector) and **Decoder**.

Encoder:

- Stack of *several* recurrent units,
- Each accepts a *single* element of the input sequence,
- Propagates information *forward*: $h_t = f(h_{t-1}, x_t; \theta)$.



Decoder:

- Stack of *several* recurrent units,
- Each predicts an output y_t
- Propagates information *forward*: $h_t = g(h_{t-1}, \zeta)$.

Context:

- Final hidden state produced from the encoder
- Goal: Encapsulate the information for all input elements

Cho, Van Merriënboer, Gulcehre, Bahdanau, Bougares, Schwenk & Bengio (2014),
Sutskever, Vinyals & Le (2014)

RNN: Three blocks of parameter/transformations:

1. Input $x_t \mapsto$ Hidden state h_t
2. Previous hidden state $h_t \mapsto$ Next hidden state h_{t+1}
3. Hidden state $h_t \mapsto$ Output y_t

RNN: Three blocks of parameter/transformations:

1. Input $x_t \mapsto$ Hidden state h_t
2. Previous hidden state $h_t \mapsto$ Next hidden state h_{t+1}
3. Hidden state $h_t \mapsto$ Output y_t

RNN architecture: Each of the blocks is associated with a single weight matrix
 \longleftrightarrow **Single layer** within a deep MLP


Deep Recurrent Networks

RNN: Three blocks of parameter/transformations:

1. Input $x_t \mapsto$ Hidden state h_t
2. Previous hidden state $h_t \mapsto$ Next hidden state h_{t+1}
3. Hidden state $h_t \mapsto$ Output y_t

RNN architecture: Each of the blocks is associated with a single weight matrix
 \longleftrightarrow **Single layer** within a deep MLP

\leadsto *We can introduce depth into each of these operations!*



Deep Learning, Goodfellow, Bengio, Courville (2016)
§ 10.5 Deep Recurrent Networks

www.deeplearningbook.org

Training Recurrent Networks

3.1 Forward Propagation


3.2 Back-Propagation Through Time

3.3 Limit of the simple RNN model

Forward Propagation (Input & output of same length)

Weight matrices: Connections between the different states

Activation functions:




Forward Propagation (Input & output of same length)

Weight matrices: Connections between the different states

1. Input-to-Hidden: U
2. Hidden-to-Hidden: W + Bias vector : b and c .
3. Hidden-to-Output: V

Activation functions:



Forward Propagation (Input & output of same length)

Weight matrices: Connections between the different states


1. Input-to-Hidden: U
2. Hidden-to-Hidden: W + Bias vector : b and c .
3. Hidden-to-Output: V

Activation functions:

1. Input-to-Hidden: **Hyperbolic tangent** activation function
2. Assume that the output is discrete.

Outputs \longleftrightarrow Unnormalized probabilities of each possible value

\leadsto **Softmax** activation function to obtain normalized probabilities



Forward Propagation (Input & output of same length)

Weight matrices: Connections between the different states


1. Input-to-Hidden: U
2. Hidden-to-Hidden: W + Bias vector : b and c .
3. Hidden-to-Output: V

Activation functions:

1. Input-to-Hidden: **Hyperbolic tangent** activation function
2. Assume that the output is discrete.

Outputs \longleftrightarrow Unnormalized probabilities of each possible value

\leadsto **Softmax** activation function to obtain normalized probabilities



Given initial state h_0 ,
For each step $t \in \llbracket 1, T \rrbracket$:

$$h_t = \tanh(W h_{t-1} + U x_t + b)$$

$$\hat{y}_t = \text{softmax}(V h_t + c)$$

Training Recurrent Networks

3.1 Forward Propagation

3.2 Back-Propagation Through Time

3.3 Limit of the simple RNN model

Computing the Gradient in a Recurrent Neural Network

Total loss function for a given sequence values x paired with a sequence of y

$$\begin{aligned}\mathcal{L}(\{x_1, \dots, x_T\}, \{y_1, \dots, y_T\}; \theta) &= \sum_{t=1}^T L_t(\theta) \quad ; \quad \theta = (U, V, W, b, c) \\ &= \sum_{t=1}^T -\log p_{\text{model}}(y_t | \{x_1, \dots, x_t\}; \theta)\end{aligned}$$

Computing the Gradient in a Recurrent Neural Network

Total loss function for a given sequence values x paired with a sequence of y

$$\begin{aligned}\mathcal{L}(\{x_1, \dots, x_T\}, \{y_1, \dots, y_T\}; \theta) &= \sum_{t=1}^T L_t(\theta) \quad ; \quad \theta = (U, V, W, b, c) \\ &= \sum_{t=1}^T -\log p_{\text{model}}(y_t | \{x_1, \dots, x_t\}; \theta)\end{aligned}$$

Gradient of this loss function **expensive operation**:

1. A **forward** propagation pass moving left to right,
2. Followed by a **backward** propagation pass moving right to left.

\Rightarrow Runtime = $O(T)$ and cannot be reduced by parallelization...

Computing the Gradient in a Recurrent Neural Network

Total loss function for a given sequence values x paired with a sequence of y

$$\begin{aligned}\mathcal{L}(\{x_1, \dots, x_T\}, \{y_1, \dots, y_T\}; \theta) &= \sum_{t=1}^T L_t(\theta) \quad ; \quad \theta = (U, V, W, b, c) \\ &= \sum_{t=1}^T -\log p_{\text{model}}(y_t | \{x_1, \dots, x_t\}; \theta)\end{aligned}$$

Gradient of this loss function **expensive operation**:

1. A **forward** propagation pass moving left to right,
2. Followed by a **backward** propagation pass moving right to left.

⇒ Runtime = $O(T)$ and cannot be reduced by parallelization...

Moreover, states computed in the forward pass must be stored until they are reused during the backward pass

⇒ Memory cost = $O(T)$

Training Recurrent Networks

3.1 Forward Propagation


3.2 Back-Propagation Through Time

3.3 Limit of the simple RNN model

Long-Term Dependencies

Strenght of RNNs: Being able to connect previous information to the present task


If the gap between relevant information and where it is needed is small, RNNs work “perfectly”




Long-Term Dependencies

Strength of RNNs: Being able to connect previous information to the present task

😊 If the gap between relevant information and where it is needed is **small**, RNNs work “perfectly”




😞 Unfortunately, as this gap **widens**, RNNs become unable to learn to connect information...



A Short Memory Model

- The **hyperbolic tangent** tends to *smash* the values taken as input.
- Information coming from x_1 passes t times through \tanh
 $\leadsto x_1$ is *very crushed*.
- At the same time, the one from x_t passes through \tanh only once
 $\leadsto x_t$ is *preserved* overall.



A Short Memory Model

- The **hyperbolic tangent** tends to *smash* the values taken as input.


- Information coming from x_1 passes t times through \tanh

$\leadsto x_1$ is very **crushed**.

- At the same time, the one from x_t passes through \tanh only once

$\leadsto x_t$ is **preserved** overall.

$x_t \gg x_1$
for predictions



A Short Memory Model

- The **hyperbolic tangent** tends to *smash* the values taken as input.


- Information coming from x_1 passes t times through \tanh

$\leadsto x_1$ is very **crushed**.

- At the same time, the one from x_t passes through \tanh only once

$\leadsto x_t$ is **preserved** overall.

$x_t \gg x_1$
for predictions




Naive approach: Compensate this crushing effect by a stronger weight associated with h_1 , via the weight matrix W .

A Short Memory Model

- The **hyperbolic tangent** tends to *smash* the values taken as input.
- Information coming from x_1 passes t times through \tanh
 $\leadsto x_1$ is very **crushed**.
- At the same time, the one from x_t passes through \tanh only once
 $\leadsto x_t$ is **preserved** overall.

$x_t \gg x_1$
for predictions



Naive approach: Compensate this crushing effect by a stronger weight associated with h_1 , via the weight matrix W .

But: Matrix W shared by all cells in the layer, i.e. all intermediate outputs weighted by the same weight.

A Short Memory Model

- The **hyperbolic tangent** tends to *smash* the values taken as input.


- Information coming from x_1 passes t times through \tanh

$\leadsto x_1$ is very **crushed**.

- At the same time, the one from x_t passes through \tanh only once

$\leadsto x_t$ is **preserved** overall.

$x_t \gg x_1$
for predictions



Naive approach: Compensate this crushing effect by a stronger weight associated with h_1 , via the weight matrix **W**.

But: Matrix **W** shared by all cells in the layer, *i.e.* all intermediate outputs weighted by the same weight.

\leadsto **LSTM, GRU, ...**

Vanishing of the Gradients

Basic problem: Gradients propagated over **many stages**

- **vanish** (most of the time)
- or **explode** (rarely, but with much damage to the optimization)

And even without it, *long-term dependencies arises from exponentially smaller weights given to long-term interactions compared to short-term ones.*

Vanishing of the Gradients

Basic problem: Gradients propagated over **many stages**

- **vanish** (most of the time)
- or **explode** (rarely, but with much damage to the optimization)

And even without it, *long-term dependencies arises from exponentially smaller weights given to long-term interactions compared to short-term ones.*

RNN \implies Composition of the same function multiple times, once per time step.
 \leadsto **Extremely nonlinear behavior**

Vanishing of the Gradients

Basic problem: Gradients propagated over **many stages**

- **vanish** (most of the time)
- or **explode** (rarely, but with much damage to the optimization)

And even without it, *long-term dependencies arises from exponentially smaller weights given to long-term interactions compared to short-term ones.*

RNN \implies Composition of the same function multiple times, once per time step.
 \leadsto **Extremely nonlinear behavior**

Toy model: No input, no activation function: $h_t = W^\top h_{t-1}$

Assume that W admits an eigen decomposition.

Vanishing of the Gradients

Basic problem: Gradients propagated over **many stages**

- **vanish** (most of the time)
- or **explode** (rarely, but with much damage to the optimization)

And even without it, *long-term dependencies arises from exponentially smaller weights given to long-term interactions compared to short-term ones.*

RNN \implies Composition of the same function multiple times, once per time step.

\rightsquigarrow **Extremely nonlinear behavior**

Toy model: No input, no activation function: $h_t = W^\top h_{t-1}$

Assume that W admits an eigen decomposition. Then

$$h_t = P^\top \Lambda^t P h_0$$

Vanishing of the Gradients

Basic problem: Gradients propagated over **many stages**

- **vanish** (most of the time)
- or **explode** (rarely, but with much damage to the optimization)

And even without it, *long-term dependencies arises from exponentially smaller weights given to long-term interactions compared to short-term ones.*

RNN \implies Composition of the same function multiple times, once per time step.
 \leadsto **Extremely nonlinear behavior**

Toy model: No input, no activation function: $h_t = W^\top h_{t-1}$

Assume that W admits an eigen decomposition. Then $h_t = P^\top \Lambda^t P h_0$

$|\lambda| < 1$ decay to zero
 $|\lambda| > 1$ explosion \implies Any component of h_0 not aligned with the largest eigenvector will eventually be discarded

Vanishing of the Gradients

Basic problem: Gradients propagated over **many stages**

- **vanish** (most of the time)
- or **explode** (rarely, but with much damage to the optimization)

And even without it, *long-term dependencies arises from exponentially smaller weights given to long-term interactions compared to short-term ones.*

RNN \implies Composition of the same function multiple times, once per time step.
 \leadsto **Extremely nonlinear behavior**


Toy model: No input, no activation function: $h_t = W^\top h_{t-1}$

Assume that W admits an eigen decomposition. Then $h_t = P^\top \Lambda^t P h_0$

$|\lambda| < 1$ decay to zero \implies Any component of h_0 not aligned with the
 $|\lambda| > 1$ explosion largest eigenvector will eventually be discarded

Remark: Problem particular to recurrent networks, due to the multiple composition of **same** function

Summary of the RNN Framework



Given initial state h_0 ,
For each step $t \in \llbracket 1, T \rrbracket$:


$$h_t = \tanh(W h_{t-1} + U x_t + b)$$

$$\hat{y}_t = \text{softmax}(V h_t + c)$$

Summary of the RNN Framework

Advantages of RNN:

- Performance not significantly affected from *missing values*;
- Can find complex patterns in the input time series;
- Good results in forecasting more than few-steps;
- Can model sequence of data so that each sample can be assumed to be dependent on previous ones.



Given initial state h_0 ,
For each step $t \in \llbracket 1, T \rrbracket$:

$$h_t = \tanh(W h_{t-1} + U x_t + b)$$

$$\hat{y}_t = \text{softmax}(V h_t + c)$$


Summary of the RNN Framework

Advantages of RNN:

- Performance not significantly affected from *missing values*;
- Can find complex patterns in the input time series;
- Good results in forecasting more than few-steps;
- Can model sequence of data so that each sample can be assumed to be dependent on previous ones.

Disadvantages of RNN:

- When trained on long time series, RNNs suffer from the **vanishing gradient** or **exploding gradient** problem;
- Weak memory unable to take into account several elements of the past in the prediction of the future;
- Training of a Recurrent Neural Network hard to parallelize and computationally expensive.



Given initial state h_0 ,
For each step $t \in \llbracket 1, T \rrbracket$:

$$h_t = \tanh(W h_{t-1} + U x_t + b)$$

$$\hat{y}_t = \text{softmax}(V h_t + c)$$

Challenge of Long-Term Dependencies

4.1 Optimization for Long-Term Dependencies

4.2 The Challenge of Long-Term Dependencies

Optimization for Long-Term Dependencies

Clipping gradient: avoids gradient explode but **NOT** gradient vanish.

- Mikolov (2012): Clip the parameter gradient from a mini-batch element-wise,
- Pascanu et al. (2013): Clip the norm of the gradient just before the parameter update.

Regularizing to encourage the information flow: Favor gradient vector being back-propagated to maintain its magnitude. (Pascanu et al., 2013)

Optimization for Long-Term Dependencies

Clipping gradient: avoids gradient explode but NOT gradient vanish.

- Mikolov (2012): Clip the parameter gradient from a mini-batch element-wise,
- Pascanu et al. (2013): Clip the norm of the gradient just before the parameter update.

Regularizing to encourage the information flow: Favor gradient vector being back-propagated to maintain its magnitude. (Pascanu et al., 2013)

Remarks: These techniques are *not* quite useful nowadays: Most of the time using LSTM will solve the long-term dependency problem.

Take-Home Message:

It is often much easier to design a model that is easy to optimize than it is to design a more powerful optimization algorithm.

Challenge of Long-Term Dependencies

4.1 Optimization for Long-Term Dependencies

4.2 The Challenge of Long-Term Dependencies

The Challenge of Long-Term Dependencies

First idea: Design a model that operates at **multiple time scales**

- (i) Some parts operate at *fine-grained* time scales
 - ↪ Handle **small details**,
- (ii) Other parts operate at *coarse* time scales
 - ↪ **Transfer information** from distant past to present

The Challenge of Long-Term Dependencies

First idea: Design a model that operates at **multiple time scales**

- (i) Some parts operate at *fine-grained* time scales
 - ↪ Handle **small details**,
- (ii) Other parts operate at *coarse* time scales
 - ↪ **Transfer information** from distant past to present


Various strategies: Skipped connections, Leaky units, Deleted connections, *etc.*

The Challenge of Long-Term Dependencies

First idea: Design a model that operates at **multiple time scales**

- (i) Some parts operate at *fine-grained* time scales
 \leadsto Handle **small details**,
- (ii) Other parts operate at *coarse* time scales
 \leadsto **Transfer information** from distant past to present

Various strategies: *Skipped connections*, Leaky units, Deleted connections, etc.



\leadsto Gradients decrease exponentially with rate $\frac{T}{d}$ rather than T .

Lin, Horne, Tino
and Giles (1996).²²

Leaky Units: Mozer (1992); El Hihi & Bengio (1996).

Removing Connections: Mozer (1992); Pascanu, Mikolov & Bengio (2013).

Leaky Units: Mozer (1992); El Hihi & Bengio (1996).

- **Idea:** Each hidden state h_t is a “summary of history”

Removing Connections: Mozer (1992); Pascanu, Mikolov & Bengio (2013).

Leaky Units: Mozer (1992); El Hihi & Bengio (1996).

- **Idea:** Each hidden state h_t is a “summary of history”, set to memorize both
 - (i) a summary of the immediate *past* h_{t-1} ,
 - (ii) and some “new stuff” of *present* time x_t .

$$h_t = \alpha h_{t-1} + (1 - \alpha) x_t$$

Removing Connections: Mozer (1992); Pascanu, Mikolov & Bengio (2013).

Strategies for Multiple Time Scales

Leaky Units: Mozer (1992); El Hihi & Bengio (1996).

- **Idea:** Each hidden state h_t is a “summary of history”, set to memorize both
 - (i) a summary of the immediate *past* h_{t-1} ,
 - (ii) and some “new stuff” of *present* time x_t .
- Parameter α substitutes to the matrix W .
 - $\alpha \simeq 1$ **remembering** information about the past for a long time;
 - $\alpha \simeq 0$ information about the past is rapidly **discarded**.

$$h_t = \alpha h_{t-1} + (1 - \alpha) x_t$$

Removing Connections: Mozer (1992); Pascanu, Mikolov & Bengio (2013).

Strategies for Multiple Time Scales

Leaky Units: Mozer (1992); El Hihi & Bengio (1996).

- **Idea:** Each hidden state h_t is a “summary of history”, set to memorize both
 - (i) a summary of the immediate *past* h_{t-1} ,
 - (ii) and some “new stuff” of *present* time x_t .

$$h_t = \alpha h_{t-1} + (1 - \alpha) x_t$$
 - Parameter α substitutes to the matrix W .
 - $\alpha \simeq 1$ **remembering** information about the past for a long time;
 - $\alpha \simeq 0$ information about the past is rapidly **discarded**.
-

Removing Connections: Mozer (1992); Pascanu, Mikolov & Bengio (2013).

- **Idea:** **Actively** removing length-one connections and replacing them with longer connections.

Strategies for Multiple Time Scales

Leaky Units: Mozer (1992); El Hihi & Bengio (1996).

- **Idea:** Each hidden state h_t is a “summary of history”, set to memorize both
 - (i) a summary of the immediate *past* h_{t-1} ,
 - (ii) and some “new stuff” of *present* time x_t .

$$h_t = \alpha h_{t-1} + (1 - \alpha) x_t$$
 - Parameter α substitutes to the matrix W .
 - $\alpha \simeq 1$ **remembering** information about the past for a long time;
 - $\alpha \simeq 0$ information about the past is rapidly **discarded**.
-

Removing Connections: Mozer (1992); Pascanu, Mikolov & Bengio (2013).

- **Idea:** **Actively** removing length-one connections and replacing them with longer connections.
- **Removing connections \neq Skipping connections:**
 - **Removing connections:** Units forced to operate on along time scale;
 - **Skip connections:** Units may learn on a long time scale but may also focus on their other, short-term connections.

Gated Recurrent Neural Network

5.1 Long-Short-Term-Memory Architecture


5.2 Variants on Long-Short-Term-Memory

5.3 Gated Recurrent Unit

Long-Short-Term-Memory Neural Networks

*"Special kind of RNN's, capable of learning **long-term** dependencies."*

Hochreiter &
Schmidhuber
(1997).




- **Idea:** divide the signal between what is important in:
 - (i) *short term* through the **hidden state** (analogous to output of a vanilla RNN),
 - (ii) *long term*, through a new state : the **cell state** (acts as a long term memory).

Long-Short-Term-Memory Neural Networks

*"Special kind of RNN's, capable of learning **long-term** dependencies."*

Hochreiter &
Schmidhuber
(1997).




- **Idea:** divide the signal between what is important in:
 - (i) *short term* through the **hidden state** (analogous to output of a vanilla RNN),
 - (ii) *long term*, through a new state : the **cell state** (acts as a long term memory).
- LSTM has a **three** step process:

Long-Short-Term-Memory Neural Networks

*"Special kind of RNN's, capable of learning **long-term** dependencies."*

Hochreiter &
Schmidhuber
(1997).



- **Idea:** divide the signal between what is important in:
 - (i) *short term* through the **hidden state** (analogous to output of a vanilla RNN),
 - (ii) *long term*, through a new state : the **cell state** (acts as a long term memory).


- LSTM has a **three** step process:

Forget gate From **cell state**, detect relevant information from the past;

Long-Short-Term-Memory Neural Networks

“Special kind of RNN’s, capable of learning *long-term* dependencies.”

Hochreiter &
Schmidhuber
(1997).



- **Idea:** divide the signal between what is important in:
 - (i) *short term* through the **hidden state** (analogous to output of a vanilla RNN),
 - (ii) *long term*, through a new state : the **cell state** (acts as a long term memory).

- LSTM has a **three** step process:


Forget gate From **cell state**, detect relevant information from the past;

Input gate From **current input**, select relevant in the long term information;

Long-Short-Term-Memory Neural Networks

“Special kind of RNN’s, capable of learning *long-term* dependencies.”

Hochreiter &
Schmidhuber
(1997).



- **Idea:** divide the signal between what is important in:
 - (i) *short term* through the **hidden state** (analogous to output of a vanilla RNN),
 - (ii) *long term*, through a new state : the **cell state** (acts as a long term memory).

- LSTM has a **three** step process:

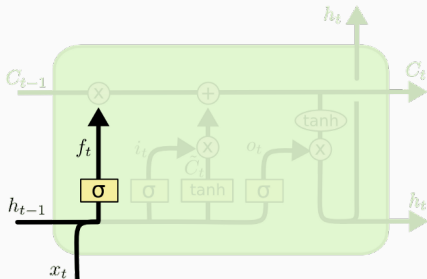
Forget gate From **cell state**, detect relevant information from the past;

Input gate From **current input**, select relevant in the long term information;

Output gate From new **cell state**, select important short term information to

Forget Gate

“Decides how much of the **past** you should remember.”



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

Input:

- Previous state h_{t-1} ,
- Content input x_t .

Layer:

- Dense layer.

Output:

- A number between 0 (omit this) and 1 (keep this) for each number in the **cell state** C_{t-1} .


↪ **Forget gate** \equiv **filter** to “forget” some information of the **cell state**.

Term to term multiplication between f_t and C_{t-1}

\implies cancellation of components of C_{t-1} whose counterparts on f_t are near 0.

Update Gate or Input Gate:

"Decides how much of *this unit* is added to the current state."



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

$$C_t = f_t \otimes C_{t-1} + i_t \otimes \tilde{C}_t$$

Input:

- Previous state h_{t-1} ,
- Content input x_t .

Output:


- Cell state C_t .

Remarks

- Input gate $i_t \equiv$ Filter,
- $\tilde{C}_t \equiv$ Candidate vector to update the cell state,
- \rightsquigarrow Update vector of the cell state.

Output Gate

“Decides which part of the **current cell** makes it to the output.”



$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t \otimes \tanh(C_t)$$

Input:

- Previous state h_{t-1} ,
- Content input x_t ,
- Cell state C_t ,

Output:

- Current state h_t .

Remarks

- Output gate $o_t \equiv$ Filter,
- New output \equiv Weighted cell filtration.


Gated Recurrent Neural Network

5.1 Long-Short-Term-Memory Architecture

5.2 Variants on Long-Short-Term-Memory

5.3 Gated Recurrent Unit

LSTM Augmented by “Peephole Connection”




- “Peephole connections”: Let the gate layers look at the cell state;
- Diagram adds peepholes to all the gates, but we can give some peepholes and not others.

Gers & Schmidhuber (2000).

$$\left\{ \begin{array}{l} f_t = \sigma(W_f \cdot [C_{t-1}, h_{t-1}, x_t] + b_f) \\ i_t = \sigma(W_i \cdot [C_{t-1}, h_{t-1}, x_t] + b_i) \\ \tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C) \\ C_t = f_t \otimes C_{t-1} + i_t \otimes \tilde{C}_t \\ o_t = \sigma(W_o \cdot [C_t, h_{t-1}, x_t] + b_o) \\ h_t = o_t \otimes \tanh(C_t) \end{array} \right.$$

Other Minor Variants

- Coupled Input and Forget gates: $f_t = 1 - i_t$
- Full gate recurrence $f_t = \sigma(W_f \cdot [h_{t-1}, x_t, C_t, f_{t-1}, i_{t-1}, o_{t-1}] + b_f)$



LSTM: A Search Space Odyssey (Greff et al., 2015)

- Tested the following variants, using Peephole LSTM as standard:

NIG No Input Gate,

NFG No Forget Gate,

NOG No Output Gate,

NIAF No Input Activation Function,

NOAF No Output Activation Function,




NP No Peepholes,

CIFG Coupled Input and Forget Gate,

FGR Full Gate Recurrence;


- On the tasks of:

- Timit Speech Recognition: Audio frame to 1 of 61 phonemes,
- IAM Online Handwriting Recognition: Sketch to characters,
- JSB Chorales: Next-step music frame prediction.



LSTM: A Search Space Odyssey (Greff et al., 2015)

- Standard LSTM performed reasonably well on multiple datasets and *none of the modifications significantly improved the performance*;
- Coupling gates and removing peephole connections simplified the LSTM without hurting performance much;
- The forget gate and output activation are crucial;
- Found interaction between learning rate and network size to be minimal – indicates calibration can be done using a small network first.



Gated Recurrent Neural Network


5.1 Long-Short-Term-Memory Architecture

5.2 Variants on Long-Short-Term-Memory

5.3 Gated Recurrent Unit

Gated Recurrent Unit


"Another special kind of RNN's, designed to *solve the vanishing gradient problem*."



Cho, et al. (2014).

Gated Recurrent Unit

"Another special kind of RNN's, designed to *solve the vanishing gradient problem*."




Cho, et al. (2014).

- A GRU unit is composed of:

Gated Recurrent Unit

“Another special kind of RNN’s, designed to *solve the vanishing gradient problem*.”




Cho, et al. (2014).

- A GRU unit is composed of:

Reset gate How much information from the previous steps can be forgotten;

Gated Recurrent Unit

“Another special kind of RNN’s, designed to *solve the vanishing gradient problem*.”



Cho, et al. (2014).


- A GRU unit is composed of:

Reset gate How much information from the previous steps can be forgotten;

Update gate How much information from the previous time steps must be saved;

Gated Recurrent Unit

“Another special kind of RNN’s, designed to *solve the vanishing gradient problem*.”



Cho, et al. (2014).


- A GRU unit is composed of:

Reset gate How much information from the previous steps can be forgotten;

Update gate How much information from the previous time steps must be saved;

Current memory Provides information throughout the sequence,
Represents the memory of the network.

Gates: Reset and Update



Reset gates:

- How to combine the new input with the previous memory?
- How much information from previous time steps can be **forgotten**?


$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t] + b_i)$$

Update gates:

- How much information from previous time steps needs to be **passed along to the future**.

$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t] + b_i)$$

Cho, Van Merriënboer, Gulcehre, Bahdanau, Bougares, Schwenk & Bengio (2014).



Current memory content:

- **Store** the relevant information from the past

$$\tilde{h}_t = \tanh(W \cdot [r_t \otimes h_{t-1}, x_t] + b_i)$$

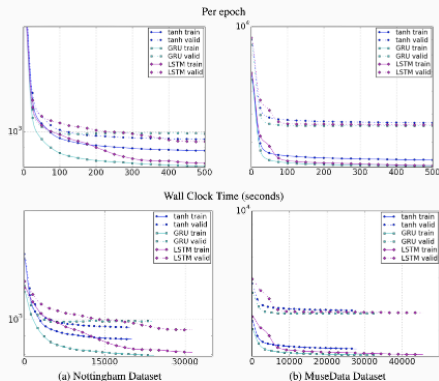
Final memory (at current time step):

- What to collect from the current memory content \tilde{h}_t and what from the previous steps h_{t-1} .

$$h_t = \sigma((1 - z_t \otimes h_{t-1} + z_t \otimes \tilde{h}_t))$$

Cho, Van Merriënboer, Gulcehre, Bahdanau, Bougares, Schwenk & Bengio (2014).

Vanilla RNN vs. LSTM vs. GRU



- GRU has fewer tensor operations
 ~ Little speedier to train than LSTM.

- No clear winner between LSTM & GRU
 ~ Try both to determine which one works better for the considered case.

			RNN	GRU	LSTM
Next step Music prediction	Nottingham	train	3.22	2.79	3.08
		test	3.13	3.23	3.20
	JSB Chorales	train	8.82	6.94	8.15
		test	9.10	8.54	8.67
	MuseData	train	5.64	5.06	5.18
		test	6.23	5.99	6.23
Speech recognition	Piano-midi	train	5.64	4.93	6.49
		test	9.03	8.82	9.03
	Ubisoft dataset A	train	6.29	2.31	1.44
		test	6.44	3.59	2.70
	Ubisoft dataset B	train	7.61	0.38	0.80
		test	7.62	0.88	1.26

Junyoung, Caglar, KyungHyun & Bengio (2014).

A Profusion of RNN Variants

- Only a **few** of the most notable LSTM variants.

Which of these variants is best? Do the differences matter?

A Profusion of RNN Variants

- Only a **few** of the most notable LSTM variants.

Which of these variants is best? Do the differences matter?

↪ Jozefowicz, et al. (2015): *Empirical exploration of RNN*

A Profusion of RNN Variants

- Only a **few** of the most notable LSTM variants.

Which of these variants is best? Do the differences matter?

↪ *Jozefowicz, et al. (2015): Empirical exploration of RNN*

- Given the rather ad-hoc design of the LSTM, the authors:
 - (i) tried to determine if the architecture of LSTM is **optimal**,

A Profusion of RNN Variants

- Only a **few** of the most notable LSTM variants.

Which of these variants is best? Do the differences matter?

↪ *Jozefowicz, et al. (2015): Empirical exploration of RNN*

- Given the rather ad-hoc design of the LSTM, the authors:
 - (i) tried to determine if the architecture of LSTM is **optimal**,
 - (ii) tested **10k** RNN architectures (including new ones),

A Profusion of RNN Variants

- Only a **few** of the most notable LSTM variants.

Which of these variants is best? Do the differences matter?

↪ *Jozefowicz, et al. (2015): Empirical exploration of RNN*

- Given the rather ad-hoc design of the LSTM, the authors:
 - (i) tried to determine if the architecture of LSTM is **optimal**,
 - (ii) tested **10k** RNN architectures (including new ones),
 - (iii) found some that worked better on certain tasks.

LSTM initialized with a large positive forget gate bias outperformed both the basic LSTM and the GRU!

A Profusion of RNN Variants

- Only a **few** of the most notable LSTM variants.

Which of these variants is best? Do the differences matter?

↪ *Jozefowicz, et al. (2015): Empirical exploration of RNN*

- Given the rather ad-hoc design of the LSTM, the authors:
 - (i) tried to determine if the architecture of LSTM is **optimal**,
 - (ii) tested **10k** RNN architectures (including new ones),
 - (iii) found some that worked better on certain tasks.

LSTM initialized with a large positive forget gate bias outperformed both the basic LSTM and the GRU!

- All architectures are evaluated on 3 problems:

Arithmetic Compute sum/difference of two numbers, with distractors


$$3e36d9 - h1h39f94eeh43keg3c = 3369 - 13994433 = -13991064;$$

XML Modeling Predict next character in valid XML modeling;

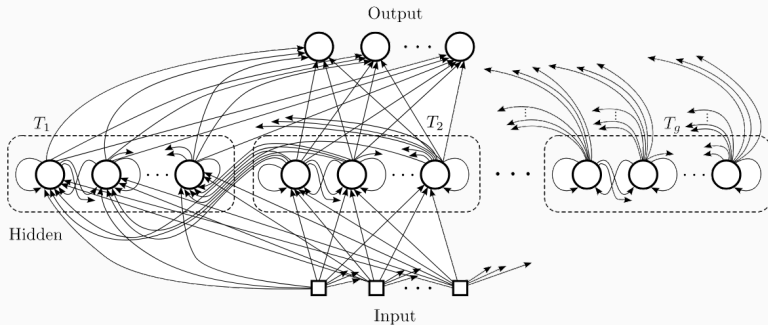
Language Modeling Predict distributions over words (Penn.

- There are also completely different strategies: Cf. **Clockwork RNN** (Koutnik, Greff, Gomez & Schmidhuber, 2014)

Clockwork Recurrent Neural Network




Clockwork Recurrent Neural Network



- “Better memory” than LSTMs: **Structured hidden layer** that does not enforce representing the mean of all inputs (running average in case of the LSTM).

Clockwork Recurrent Neural Network



- “Better memory” than LSTMs: **Structured hidden layer** that does not enforce representing the mean of all inputs (running average in case of the LSTM).
- *But not really used in practice:* Need to know the **numbers of modules** to partition the hidden layer,
 - Depends most likely on the data,
 - No literature that suggests how to do this properly.

Explicit Memory

6.1 Memory Networks

6.2 Neural Turing Machines

Memory Networks : Pushing Further the Idea of Memory in Neural Networks

- Neural networks excel at storing *implicit knowledge*.
- However, they struggle to memorize *facts*.

Memory Networks : Pushing Further the Idea of Memory in Neural Networks

- Neural networks excel at storing *implicit knowledge*.
- However, they struggle to memorize *facts*.

hypothesize: (Graves et al., 2014) *Neural nets lack of working memory:*

- Stochastic gradient descent requires **many** presentations of the **same** input before it can be stored in neural network parameters,
- And even then, that input will not be stored especially precisely.

Memory Networks : Pushing Further the Idea of Memory in Neural Networks

- Neural networks excel at storing *implicit knowledge*.
- However, they struggle to memorize *facts*.

hypothesize: (Graves et al., 2014) *Neural nets lack of **working memory**:*

- Stochastic gradient descent requires **many** presentations of the **same** input before it can be stored in neural network parameters,
- And even then, that input will not be stored especially precisely.

~> *Introduction of an **explicit memory** component.*

1. **Memory networks** (Weston, Chopra & Bordes, 2014)
 - Set of memory cells that can be accessed via an addressing mechanism,
 - *But* requires a supervision signal instructing them how to use their memory cells.
2. **Neural Turing Machine (NTM)** (Graves, Wayne & Danihelka, 2014) : able to learn to *read* from and *write* arbitrary content to memory cells without explicit supervision about which actions to undertake.

Memory Networks : Pushing Further the Idea of Memory in Neural Networks

- Neural networks excel at storing *implicit knowledge*.
- However, they struggle to memorize *facts*.

hypothesize: (Graves et al., 2014) *Neural nets lack of working memory:*

- Stochastic gradient descent requires **many** presentations of the **same** input before it can be stored in neural network parameters,
- And even then, that input will not be stored especially precisely.

~> *Introduction of an **explicit memory** component.*

1. **Memory networks** (Weston, Chopra & Bordes, 2014)
 - Set of memory cells that can be accessed via an addressing mechanism,
 - *But* requires a supervision signal instructing them how to use their memory cells.
2. **Neural Turing Machine (NTM)** (Graves, Wayne & Danihelka, 2014) : able to learn to *read* from and *write* arbitrary content to memory cells without explicit supervision about which actions to undertake.

- **Memory cell in NTM** \longleftrightarrow *Memory cells in LSTMs and GRUs.*

NTM outputs an internal state that chooses which cell to read from or write to, (think of memory accesses in a computer read from or write to a specific address).


Explicit Memory

6.1 Memory Networks

6.2 Neural Turing Machines

Neural Turing Machines


- Designed such that every component of the architecture is **differentiable**, making it straight forward to train with gradient descent.



Graves et al. (2014). See also <https://distill.pub/2016/augmented-rnns/>

Neural Turing Machines


- Designed such that every component of the architecture is **differentiable**, making it straight forward to train with gradient descent.
- Controller**: LSTM or Feedforward neural net.



Graves et al. (2014). See also <https://distill.pub/2016/augmented-rnns/>

Neural Turing Machines


- Designed such that every component of the architecture is **differentiable**, making it straight forward to train with gradient descent.
- Controller**: LSTM or Feedforward neural net.
- Reading** head: Convex combination of the row-vectors M_t in memory: $r_t = W_t \otimes M_t$.



Graves et al. (2014). See also <https://distill.pub/2016/augmented-rnns/>

Neural Turing Machines

- Designed such that every component of the architecture is **differentiable**, making it straight forward to train with gradient descent.
- Controller**: LSTM or Feedforward neural net.
- Reading** head: Convex combination of the row-vectors M_t in memory: $r_t = W_t \otimes M_t$.
- Writing** head: Two parts:
 - Erase** \leftrightarrow Forget gate of LSTM. Given an *erase* vector, e_t : $\tilde{M}_t = M_{t-1} (1 - W_t e_t)$,
 - Add** \leftrightarrow Input gate of LSTM. Given an *add* vector, a_t : $M_t = \tilde{M}_t + W_t a_t$.



Graves et al. (2014). See also <https://distill.pub/2016/augmented-rnns/>

Even Longer-Term Memory

Experiments: Copy a sequence of binary vectors, of size between 1 and 20.

Focus:

- Can NTM store and recall a **long** sequence of arbitrary information?
- And, in particular, is it able to overcome **longer delays** than LSTMs?

Even Longer-Term Memory

Experiments: Copy a sequence of binary vectors, of size between 1 and 20.

Focus:

- Can NTM store and recall a **long** sequence of arbitrary information?
- And, in particular, is it able to overcome **longer delays** than LSTMs?




Figure 3: Copy Learning Curves.

Graves, Wayne & Danihelka (2014)




Figure 4: NTM Generalisation on the Copy Task.




Figure 5: LSTM Generalisation on the Copy Task.

But: How to know *where* to read and write? \rightsquigarrow **Attention mechanism**

Appendix: Categorical Variable Encoding

Encoding: Convert a categorical variable to numerical values for machine learning model building.

Different types of Encoding:

Micci-Barreca (2001)

■ **Label Encoding**

■ **One Hot Encoding**

■ **Target Encoding**

Feature	Feature
A	A
B	B
C	C
B	B
C	C
A	A
A	A
C	C
B	B
B	B
C	C
A	A

Target
0.39
0.24
2.21
0.31
0.76
-0.74
0.27
4.01
2.28
0.19
2.03
-0.05

Feature
A
B
C
B
C
A
A
C
B
B
C
A

Credit: Brendan Hasz blog post

Appendix: Categorical Variable Encoding

Encoding: Convert a categorical variable to numerical values for machine learning model building.

Different types of Encoding:

Micci-Barreca (2001)

■ Label Encoding

■ One Hot Encoding

■ Target Encoding

Feature
0
1
2
1
2
0
0
2
1
1
2
0

Feature
A
B
C
B
C
A
A
C
B
B
C
A

Target
0.39
0.24
2.21
0.31
0.76
-0.74
0.27
4.01
2.28
0.19
2.03
-0.05

Feature
A
B
C
B
C
A
A
C
B
B
C
A

Credit: Brendan Hasz blog post

Appendix: Categorical Variable Encoding

Encoding: Convert a categorical variable to numerical values for machine learning model building.

Different types of Encoding:

Micci-Barreca (2001)

■ Label Encoding

■ One Hot Encoding

■ Target Encoding

Feature
0
1
2
1
2
0
0
2
1
1
2
0

Feature_A	Feature_B	Feature_C
1	0	0
0	1	0
0	0	1
0	1	0
0	0	1
1	0	0
1	0	0
0	0	1
0	1	0
0	1	0
0	0	1
1	0	0

Target
0.39
0.24
2.21
0.31
0.76
-0.74
0.27
4.01
2.28
0.19
2.03
-0.05

Feature
A
B
C
B
C
A
A
C
B
B
C
A

Credit: Brendan Hasz blog post

Appendix: Categorical Variable Encoding

Encoding: Convert a categorical variable to numerical values for machine learning model building.

Different types of Encoding:

Micci-Barreca (2001)

- Label Encoding
- One Hot Encoding
- Target Encoding

Feature
0
1
2
1
2
0
0
2
1
1
2
0

Feature_A	Feature_B	Feature_C
1	0	0
0	1	0
0	0	1
0	1	0
0	0	1
1	0	0
1	0	0
0	0	1
0	1	0
0	1	0
0	0	1
1	0	0

Target
0.39
0.24
2.21
0.31
0.76
-0.74
0.27
4.01
2.28
0.19
2.03
-0.05

Feature
-0.03
0.76
2.25
0.76
2.25
-0.03
-0.03
2.25
0.76
0.76
2.25
-0.03

Credit: Brendan Hasz blog post

- Bahdanau, D., Cho, K., and Bengio, Y. (2014). **Neural machine translation by jointly learning to align and translate.** [arXiv preprint arXiv:1409.0473](#).
- Bengio, Y., Simard, P., and Frasconi, P. (1994). **Learning long-term dependencies with gradient descent is difficult.** [IEEE transactions on neural networks](#), 5(2):157–166.
- Cho, K., Van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., and Bengio, Y. (2014). **Learning phrase representations using rnn encoder-decoder for statistical machine translation.** [arXiv preprint arXiv:1406.1078](#).
- Chung, J., Gulcehre, C., Cho, K., and Bengio, Y. (2014). **Empirical evaluation of gated recurrent neural networks on sequence modeling.** [arXiv preprint arXiv:1412.3555](#).
- El Hihi, S. and Bengio, Y. (1995). **Hierarchical recurrent neural networks for long-term dependencies.** In [Nips](#), volume 409.
- Elman, J. L. (1990). **Finding structure in time.** [Cognitive science](#), 14(2):179–211.
- Gers, F. A. and Schmidhuber, J. (2000). **Recurrent nets that time and count.** In [Proceedings of the IEEE-INNS-ENNS International Joint Conference on Neural Networks. IJCNN 2000. Neural Computing: New Challenges and Perspectives for the New Millennium](#), volume 3, pages 189–194. IEEE.
- Goodfellow, I., Bengio, Y., and Courville, A. (2016). **Deep learning.** MIT press.
- Graves, A. (2012). **Supervised sequence labelling.** In [Supervised sequence labelling with recurrent neural networks](#), pages 5–13. Springer.

- Graves, A., Wayne, G., and Danihelka, I. (2014). **Neural turing machines.** arXiv preprint arXiv:1410.5401.
- Hochreiter, S. (1991). **Untersuchungen zu dynamischen neuronalen netzen.** Diploma, Technische Universität München, 91(1).
- Hochreiter, S. and Schmidhuber, J. (1997). **Long short-term memory.** Neural computation, 9(8):1735–1780.
- Hopfield, J. J. (1982). **Neural networks and physical systems with emergent collective computational abilities.** Proceedings of the national academy of sciences, 79(8):2554–2558.
- Jozefowicz, R., Zaremba, W., and Sutskever, I. (2015). **An empirical exploration of recurrent network architectures.** In International conference on machine learning, pages 2342–2350. PMLR.
- Koutnik, J., Greff, K., Gomez, F., and Schmidhuber, J. (2014). **A clockwork rnn.** In International Conference on Machine Learning, pages 1863–1871. PMLR.
- Lin, T., Horne, B. G., Tino, P., and Giles, C. L. (1998). **Learning long-term dependencies is not as difficult with narx recurrent neural networks.** Technical report.
- Little, W. A. (1974). **The existence of persistent states in the brain.** Mathematical biosciences, 19(1-2):101–120.
- Maier, A. (2020). **Lecture notes on Deep Learning.** Friedrich-Alexander-Universität Erlangen-Nürnberg (FAU).
- Micci-Barreca, D. (2001). **A preprocessing scheme for high-cardinality categorical attributes in classification and prediction problems.** ACM SIGKDD Explorations Newsletter, 3(1):27–32.
- Mikolov, T. et al. (2012). **Statistical language models based on neural networks.** Presentation at Google, Mountain View, 2nd April, 80:26.

- Mozer, M. C. (1991). **Induction of multiscale temporal structure.** Advances in neural information processing systems, 4.
- Olah, C. **Understanding LSTM networks.** <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>.
- Olah, C. and Carter, S. **Attention and augmented recurrent neural networks.** <https://distill.pub/2016/augmented-rnns/>.
- Pascanu, R., Mikolov, T., and Bengio, Y. (2013). **On the difficulty of training recurrent neural networks.** In International conference on machine learning, pages 1310–1318. PMLR.
- Rumelhart, D. E., Hinton, G. E., and Williams, R. J. (1986). **Learning representations by back-propagating errors.** nature, 323(6088):533–536.
- Schuster, M. and Paliwal, K. K. (1997). **Bidirectional recurrent neural networks.** IEEE transactions on Signal Processing, 45(11):2673–2681.
- Siegelmann, H. T. and Sontag, E. D. (1995). **On the computational power of neural nets.** Journal of computer and system sciences, 50(1):132–150.
- Sutskever, I., Vinyals, O., and Le, Q. V. (2014). **Sequence to sequence learning with neural networks.** In Advances in neural information processing systems, pages 3104–3112.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). **Attention is all you need.** Advances in neural information processing systems, 30.
- Weston, J., Chopra, S., and Bordes, A. (2014). **Memory networks.** arXiv preprint arXiv:1410.3916.