

# Machine learning under physical constraints

## Introduction to RNN

Sixin Zhang  
([sixin.zhang@toulouse-inp.fr](mailto:sixin.zhang@toulouse-inp.fr))

# Outline

Recurrent Neural Networks (RNN)

Training strategies of RNN

# Outline

Recurrent Neural Networks (RNN)

Training strategies of RNN

# RNN for sequence processing

- ▶ Many data such as time-series, language, speech, genomics can be represented in a form of sequence:  $x_1, x_2, \dots$ .
- ▶ How to process such data of various type and length?
- ▶ Two basic ideas: recurrent and convolutional.
- ▶ Goal of this lecture: define what is RNN, and to construct, train and use it.
- ▶ Reference: Ian Goodfellow, Yoshua Bengio, Aaron Courville. Deep Learning. MIT Press, 2016

# What is Recurrent?

- ▶ View 1: Output of a dynamical system is fed back to some of its inputs, e.g. time-delayed system.

$$\frac{d}{dt}x(t) = f(t, x(t), x(t - \tau)), \quad \tau > 0$$

- ▶ View 2: Finite-state automata (or Turing machine)
  - ▶ Change from one state to another in response to some inputs
  - ▶ Computers are recurrent !

## RNN: dynamical system view

- ▶ Assume state of a dynamical system at time  $t$  is  $h_t$
- ▶ Classical form ( $\theta$  is a parameter)

$$h_t = f(h_{t-1}; \theta)$$


- ▶ Input (external-signal) dependent form

$$h_t = f(h_{t-1}, x_t; \theta)$$

- ▶  $h_t$  can be interpreted as hidden states in NN.

## RNN: Automata view

- Automata: an abstract machine that can be in exactly one of a finite number of states at any given time



**Figure:** Representation of an acceptor; this example shows one that determines whether a binary number has an even number of 0s, where S1 is an accepting state and S2 is a non accepting state. See [https://en.wikipedia.org/wiki/Finite-state\\_machine](https://en.wikipedia.org/wiki/Finite-state_machine)

# Unrolling a dynamical system

- ▶  $h_t$  depends on  $x_t, x_{t-1}, \dots, x_1$  and  $h_0$ .
- ▶ To see this, unroll recursively the hidden states:

$$h_t = f(h_{t-1}, x_t; \theta) = f(f(h_{t-2}, x_{t-1}; \theta), x_t; \theta)$$





Figure: Unrolled computational graph



# Definition of RNN

- ▶ RNN: a family of NN constructed from the idea of unrolling.
- ▶ Several examples:



**Figure:** One hidden-layer RNN: allow the network's hidden units to see its (own) previous output

## Example 1

- ▶ A concrete case of one hidden-layer RNN:

$$a_t = Ux_t + b + Wh_{t-1}$$


$$h_t = \tanh(a_t)$$

$$o_t = c + Vh_t$$

$$\hat{y}_t = \text{softmax}(o_t)$$

- ▶ Input-to-hidden parameters:  $U, b$
- ▶ Hidden-to-hidden parameters:  $W, b$
- ▶ Hidden-to-output parameters:  $V, c$
- ▶ A loss  $L_t$  is then computed based on  $y_t$  and  $\hat{y}_t$ .

## Example 2



**Figure:** Allow the network's hidden units to see the previous output

## Example 3




Figure: Feedback only come from the final output

## Example 3

- ▶ A concrete case of TP
  - ▶ Goal: learn an optimal output  $o_T$  close to the target  $y_T$ .
  - ▶ Initial input  $x_0 \sim p = \mathcal{N}(\mu, \sigma^2)$
  - ▶  $h_0 = x_0$
  - ▶  $h_t = h_{t-1} + \theta$
  - ▶  $o_T = h_T$
  - ▶  $y_T = \mu$
  - ▶  $L_T = (y_T - o_T)^2$
  - ▶ Optimization problem:  $\min_{\theta \in \mathbb{R}} \mathbb{E}_{x_0 \sim p}(L_T)$

# Elman RNN

- ▶ Finding structure in time (Elman 1990)
- ▶ Represent **time** by the effect it has on processing.
- ▶ Applications: sequential prediction, language understanding.




Figure: Save effects of time in a context state, aka memory

# Example of Elman RNN: Data assimilation networks




Figure: Unroll Elman RNN over time

- ▶ Relation with Elman RNN and Data assimilation networks
  - ▶ Hidden  $h_t$ : analysis posterior distribution
  - ▶ Context  $c_t$ : prediction posterior distribution
  - ▶ Input  $x_t$ : observed state in ODS

# Advanced RNN

- ▶ Next-character generation (basis of **Chat-GPT**): **demo**
- ▶ Multiple hidden-layer (deep) RNN: audio **source separation** (demo)
- ▶ Bi-directional RNN: language model (read a sequence in 2 directions)
- ▶ Auto-encoder RNN: sequence to sequence language **translation**
- ▶ GRU, LSTM: long-dependency in signal (**RNNnoise**: demo)



# Outline

Recurrent Neural Networks (RNN)

Training strategies of RNN

# Back-propagation over time (BPTT)

- ▶ To train a RNN, one often uses the gradient of the parameters for efficient optimization. BPTT is a way to compute such gradients.
- ▶ It is based on the same idea of back-propagation in neural networks, but it is subtle due to the shared parameters across time.
- ▶ Nowadays, one can use automatic differentiation to do this. But one still needs to understand it to go further.

# Example 1

- ▶ The concrete case

$$a_t = Ux_t + b + Wh_{t-1}$$

$$h_t = \tanh(a_t)$$

$$o_t = c + Vh_t$$

$$\hat{y}_t = \text{softmax}(o_t)$$

- ▶ What is the gradient of the softmax layer,  $\hat{y}_t = \text{softmax}(o_t)$ ?
- ▶ A loss  $L_t$  is computed based on  $y_t$  and  $\hat{y}_t$ .
- ▶ How to compute the total loss  $L = \sum_t L_t$  with respect to all the parameters?

## Example 1: BPTT by chain rule

- ▶ Total loss  $L = \sum_t L_t$

$$\frac{\partial L}{\partial L_t} = 1$$

- ▶ From  $L_t$  to  $o_t$

$$\nabla_{o_t} L = \left( \frac{\partial L}{\partial o_t} \right)^T = \left( \frac{\partial L}{\partial L_t} \frac{\partial L_t}{\partial o_t} \right)^T = \left( \frac{\partial L_t}{\partial o_t} \right)^T$$

- ▶ From  $o_t$  to  $h_t$ : output  $o_t = c + Vh_t$

$$\nabla_{h_t} L = V^T \nabla_{o_t} L$$

## Example 1: BPTT by chain rule

- ▶ Compute gradient of  $h_t$  from  $h_{t+1}$ .
- ▶ Hidden states:

$$h_{t+1} = \tanh(Ux_{t+1} + b + Wh_t)$$

- ▶ Recursive relation

$$\nabla_{h_t} L = \left( \frac{\partial h_{t+1}}{\partial h_t} \right)^T \nabla_{h_{t+1}} L + \left( \frac{\partial o_t}{\partial h_t} \right)^T \nabla_{o_t} L$$

## Example 1: BPTT by chain rule

- ▶ How about the parameters, e.g.  $(U, b, W)$ ?

$$h_t = \tanh(Ux_t + b + Wh_{t-1})$$

- ▶ Idea: **accumulate all the gradients over  $t$ .**
- ▶ Write  $b$  as  $b_t$  to be clear of the partial derivatives:

$$\nabla_b L = \sum_t \left( \frac{\partial h_t}{\partial b_t} \right)^T \nabla_{h_t} L$$

where  $h_t = \tanh(Ux_t + b_t + Wh_{t-1})$ .

# Deterministic (BPTT) and online (truncated BPTT)

- ▶ Initialize  $\theta = \theta^{(0)}$
- ▶ Deterministic optimizer (BPTT) at each iteration  $k$ :

$$\theta^{(k+1)} = \theta^{(k)} - \eta_k \nabla_{\theta} L(\theta^{(k)})$$

- ▶ Online optimizer (truncated BPTT) at iteration  $k$ :

$$\theta^{(k+1)} = \theta^{(k)} - \eta_k \tilde{\nabla}_{\theta} L_{k+1}(\theta^{(k)})$$

## Example 1: Truncated BPTT

- ▶ The cost (both CPU and memory) to compute  $\nabla_b L$  is  $O(T)$  due to the summation over  $t \leq T$ . This is prohibitive when  $T$  is very large.
- ▶ Truncated BPTT reduces this cost by focusing on the impact of the “current” parameter  $b_t$  on the current loss  $L_t$ .
- ▶ The truncated gradient of  $b$  at time  $t$  is

$$\tilde{\nabla}_b L_t = \left( \frac{\partial h_t}{\partial b_t} \right)^T \nabla_{h_t} L_t, \quad \nabla_{h_t} L_t = \left( \frac{\partial o_t}{\partial h_t} \right)^T \nabla_{o_t} L_t$$

- ▶ The cost to compute  $\tilde{\nabla}_b L_t$  is  $O(1)$ .
- ▶ Discussion of  $p$ -truncated BPTT ( $p = 1, 2, \dots$ ) in the paper: Corentin Tallec, Yann Ollivier. Unbiasing Truncated Backpropagation Through Time.



# Updates of truncated BPTT

- ▶ **Principle:** Step  $k$  is represented by  $(\theta^{(k)}, \tilde{h}_k)$ . It is updated to  $(\theta^{(k+1)}, \tilde{h}_{k+1})$  using the loss at time  $k+1$ ,

$$\begin{aligned}\theta^{(k+1)} &= \theta^{(k)} - \eta_k \tilde{\nabla}_{\theta} L_{k+1}(\theta^{(k)}), \\ \tilde{h}_{k+1} &= f(\tilde{h}_k, x_{k+1}; \theta^{(k)}),\end{aligned}$$

where **the truncated gradient**

$$\tilde{\nabla}_{\theta} L_{k+1}(\theta^{(k)}) = \nabla_{\theta} \ell(f(\tilde{h}_k, x_{k+1}; \theta), y_{k+1})|_{\theta=\theta_k}.$$

- ▶ **Difference to BPTT:** In BPTT, the step  $k$  is represented by  $\theta^{(k)}$  and updated to  $\theta^{(k+1)}$  using the loss over time  $t \leq T$ .

## Example 4: BPTT vs. Truncated BPTT

- ▶ Consider the following problem:
  - ▶  $h_0 = x_0 \sim \mathcal{N}(\mu, \sigma^2)$
  - ▶  $h_t = f(h_{t-1}; \theta) = h_{t-1} + \theta$
  - ▶  $L_t = \ell(h_t) = \mathbb{E}_{x_0 \sim p}(h_t - \mu)^2$
  - ▶  $\min_{\theta \in \mathbb{R}} L = \sum_{t=1}^T L_t$
- ▶ What is the optimal solution of  $L$ ?
- ▶ What is the gradient  $\nabla_{\theta} L$  and the truncated gradient  $\tilde{\nabla}_{\theta} L_t$ ?
- ▶ (Exercise) What is the BPTT and Truncated BPTT algorithm for this problem.
- ▶ (Exercise) How to choose the learning rate  $\eta_k$  to find an optimal  $\theta$ ?